

Monitoring a production-ready microservices


by Stanislav Kolenkin, Senior DevOps.

Introduction


Microservices are solving many architectural problems that have plagued developers and operations for years. But while the benefits are now clear, many professionals aren't fully aware of the many challenges that come with microservices. A significant portion of those center around monitoring and observability.

Microservices add complexity to your system, communication, and network, which makes monitoring more complicated.

Introduction



Monolithic Architecture



Microservice Architecture

Introduction



You have a lot more systems

Introduction

50+ microservices

3 environments (DEV, STAGE, PROD)

3 replica for each service


20+ check per service

...

Think about monitoring from the start



Monitoring



A network cable is unplugged

Monitoring

Kubernetes tools:

- Kubernetes UI
- cAdvisor
- Heapster
- Kubernetes dashboard
- Grafana
-



Monitoring

What network activity
is my namespace
doing?

What transactions is my
Redis service serving?

Is my X-I-Am-Awesome
HTTP header being set?




How the hell does
my K8S service
work?!

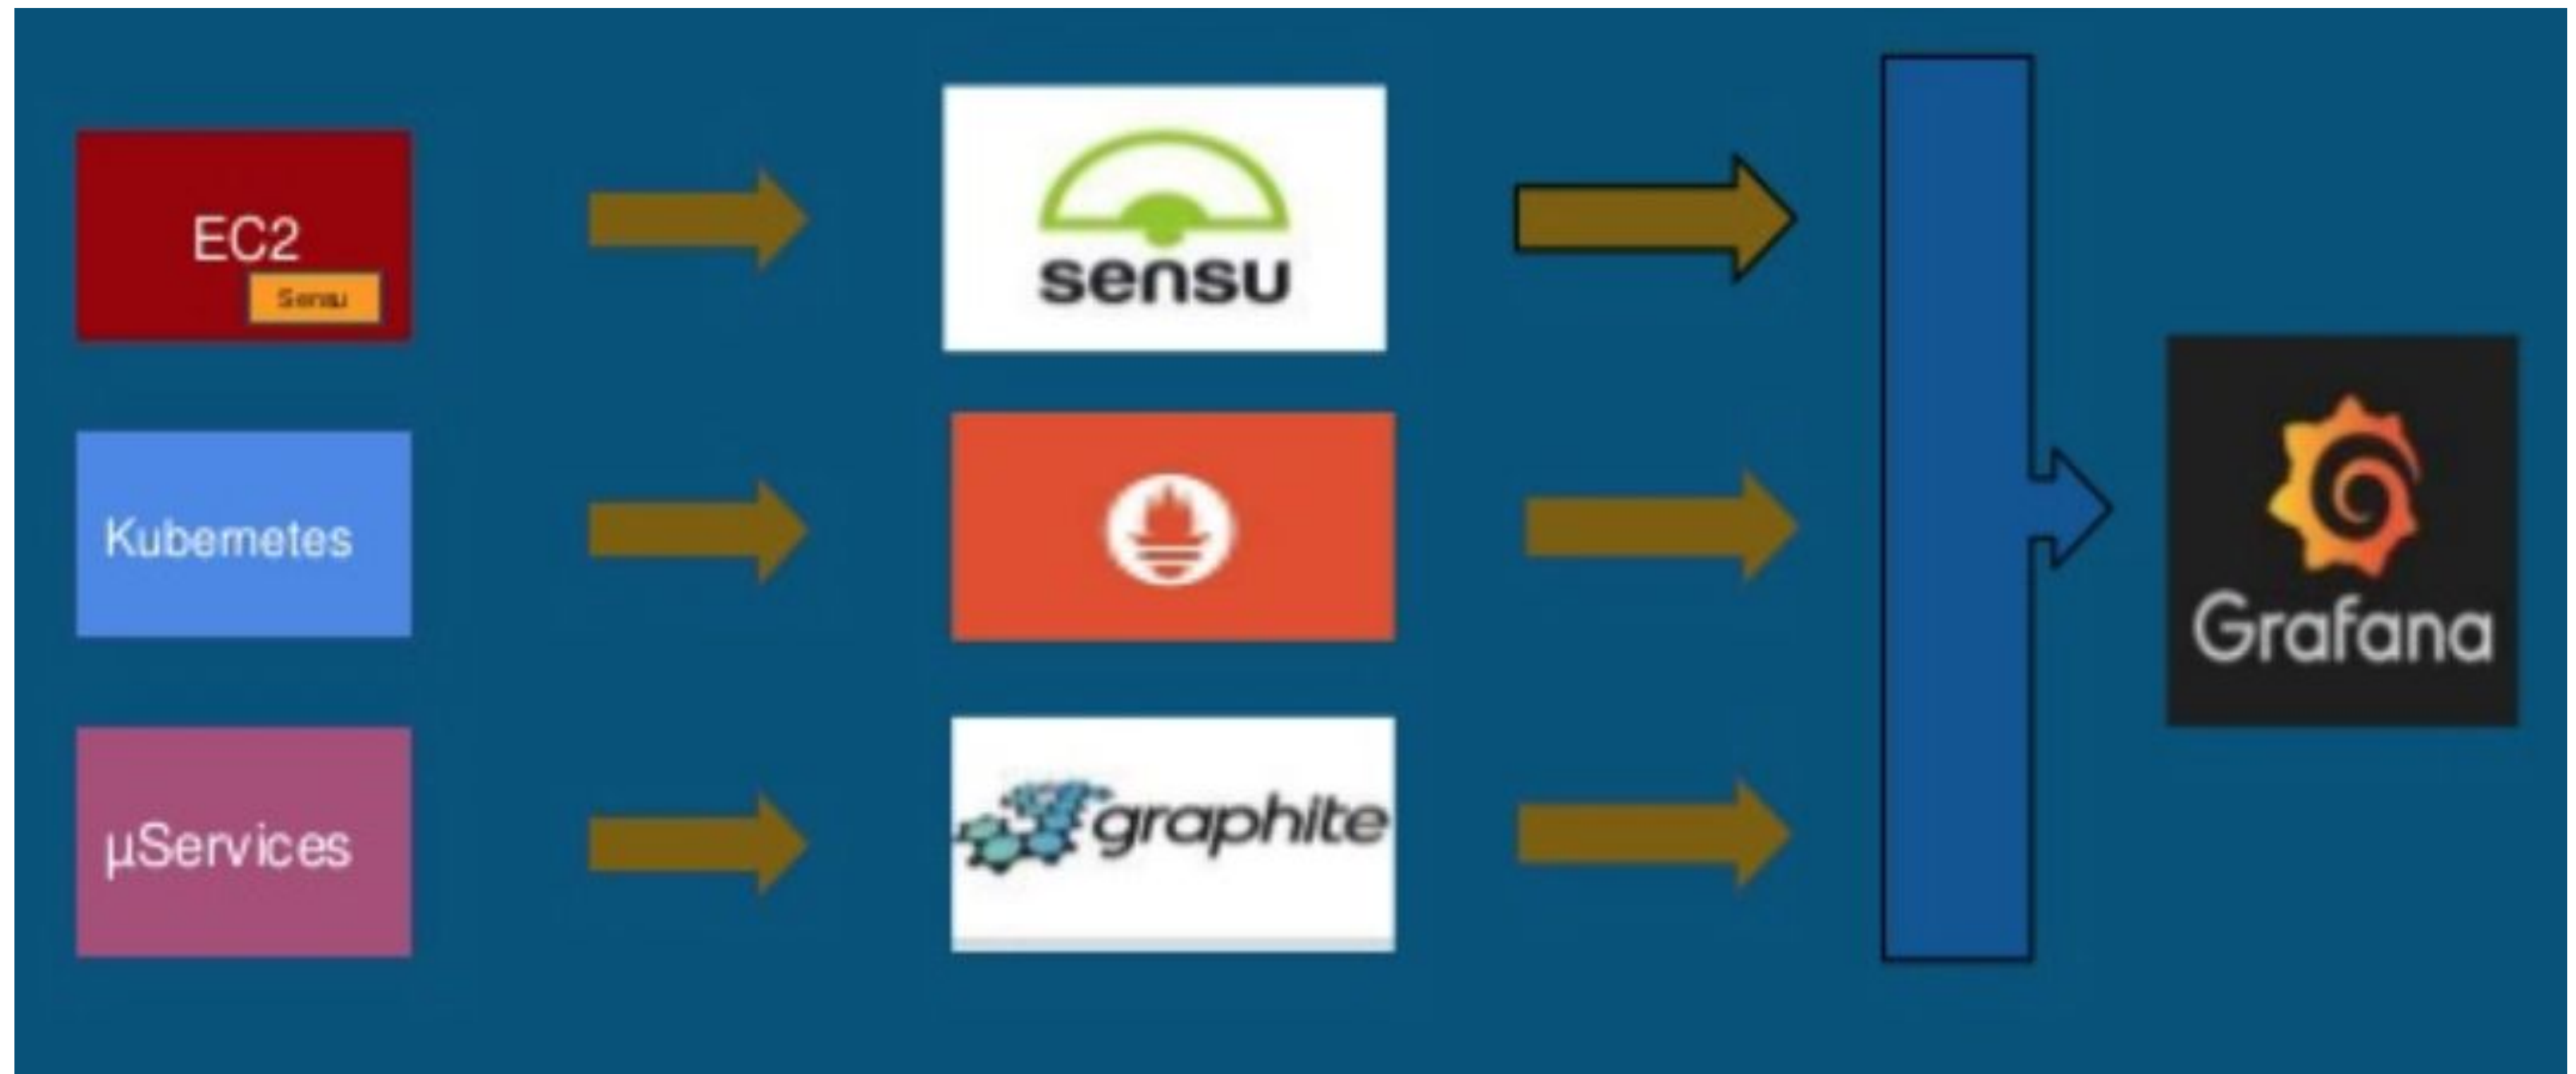
What processes are
being run in this
pod?

What files are
being written to
the file system?

Monitoring



Monitoring



Monitoring

Cons:

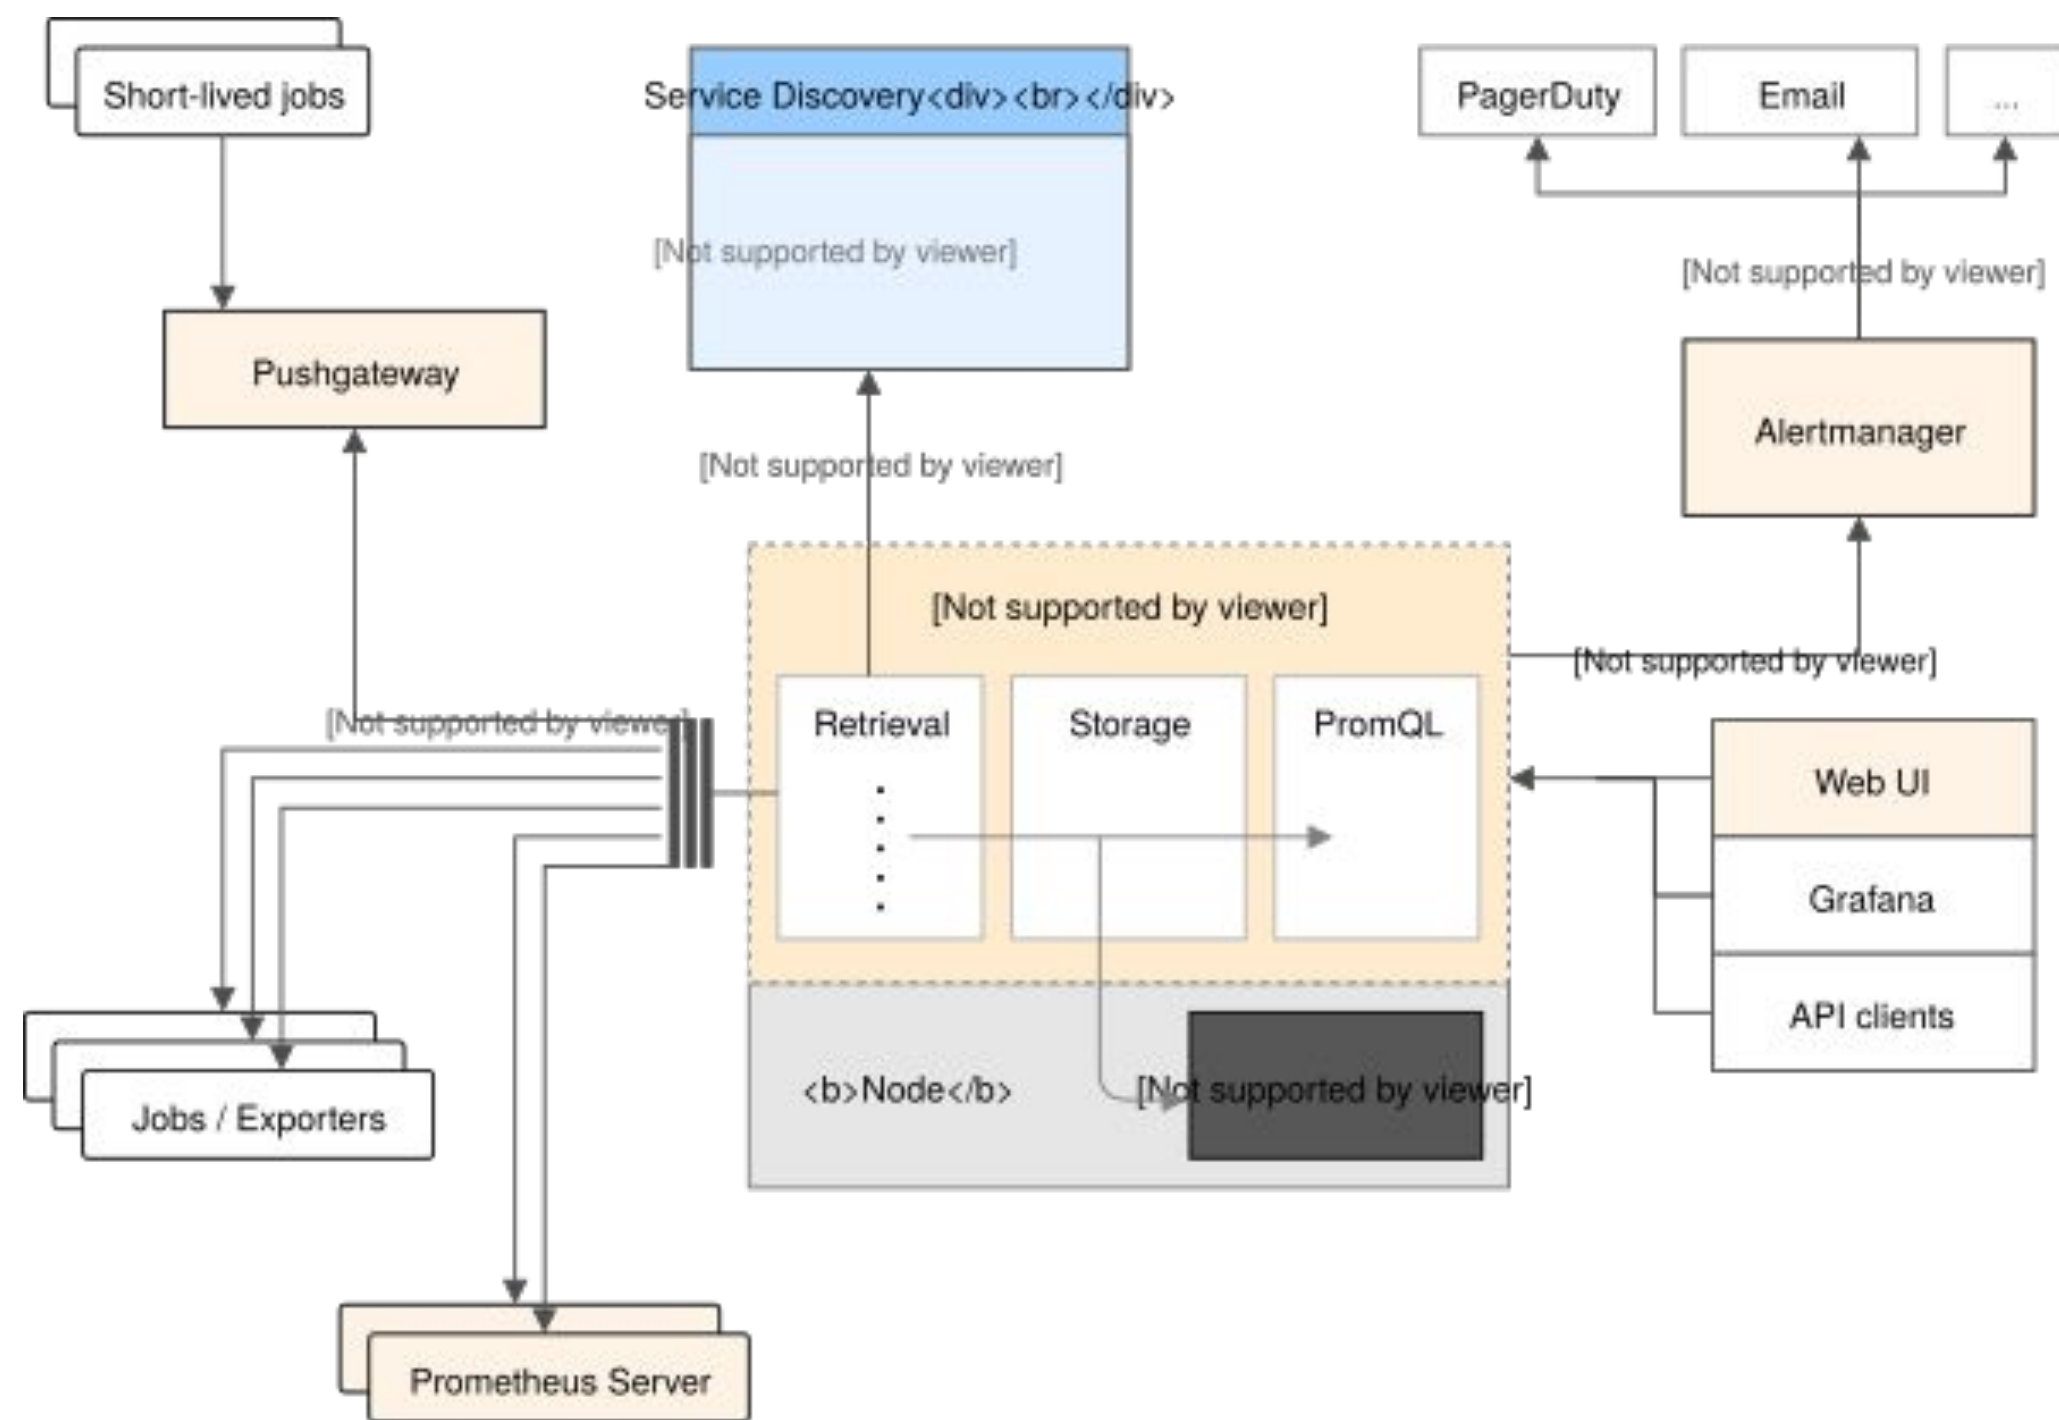
1. Multiple monitoring systems
2. Difficulty in troubleshooting
3. Additional Infrastructure cost to support three monitoring system
4. Graphite doesn't provide pod level Application metrics
5. Infra team need to understand Sensu, Prometheus alerting
6. Application metrics are single dimension
7. Grafana alerting for Application metrics

Monitoring

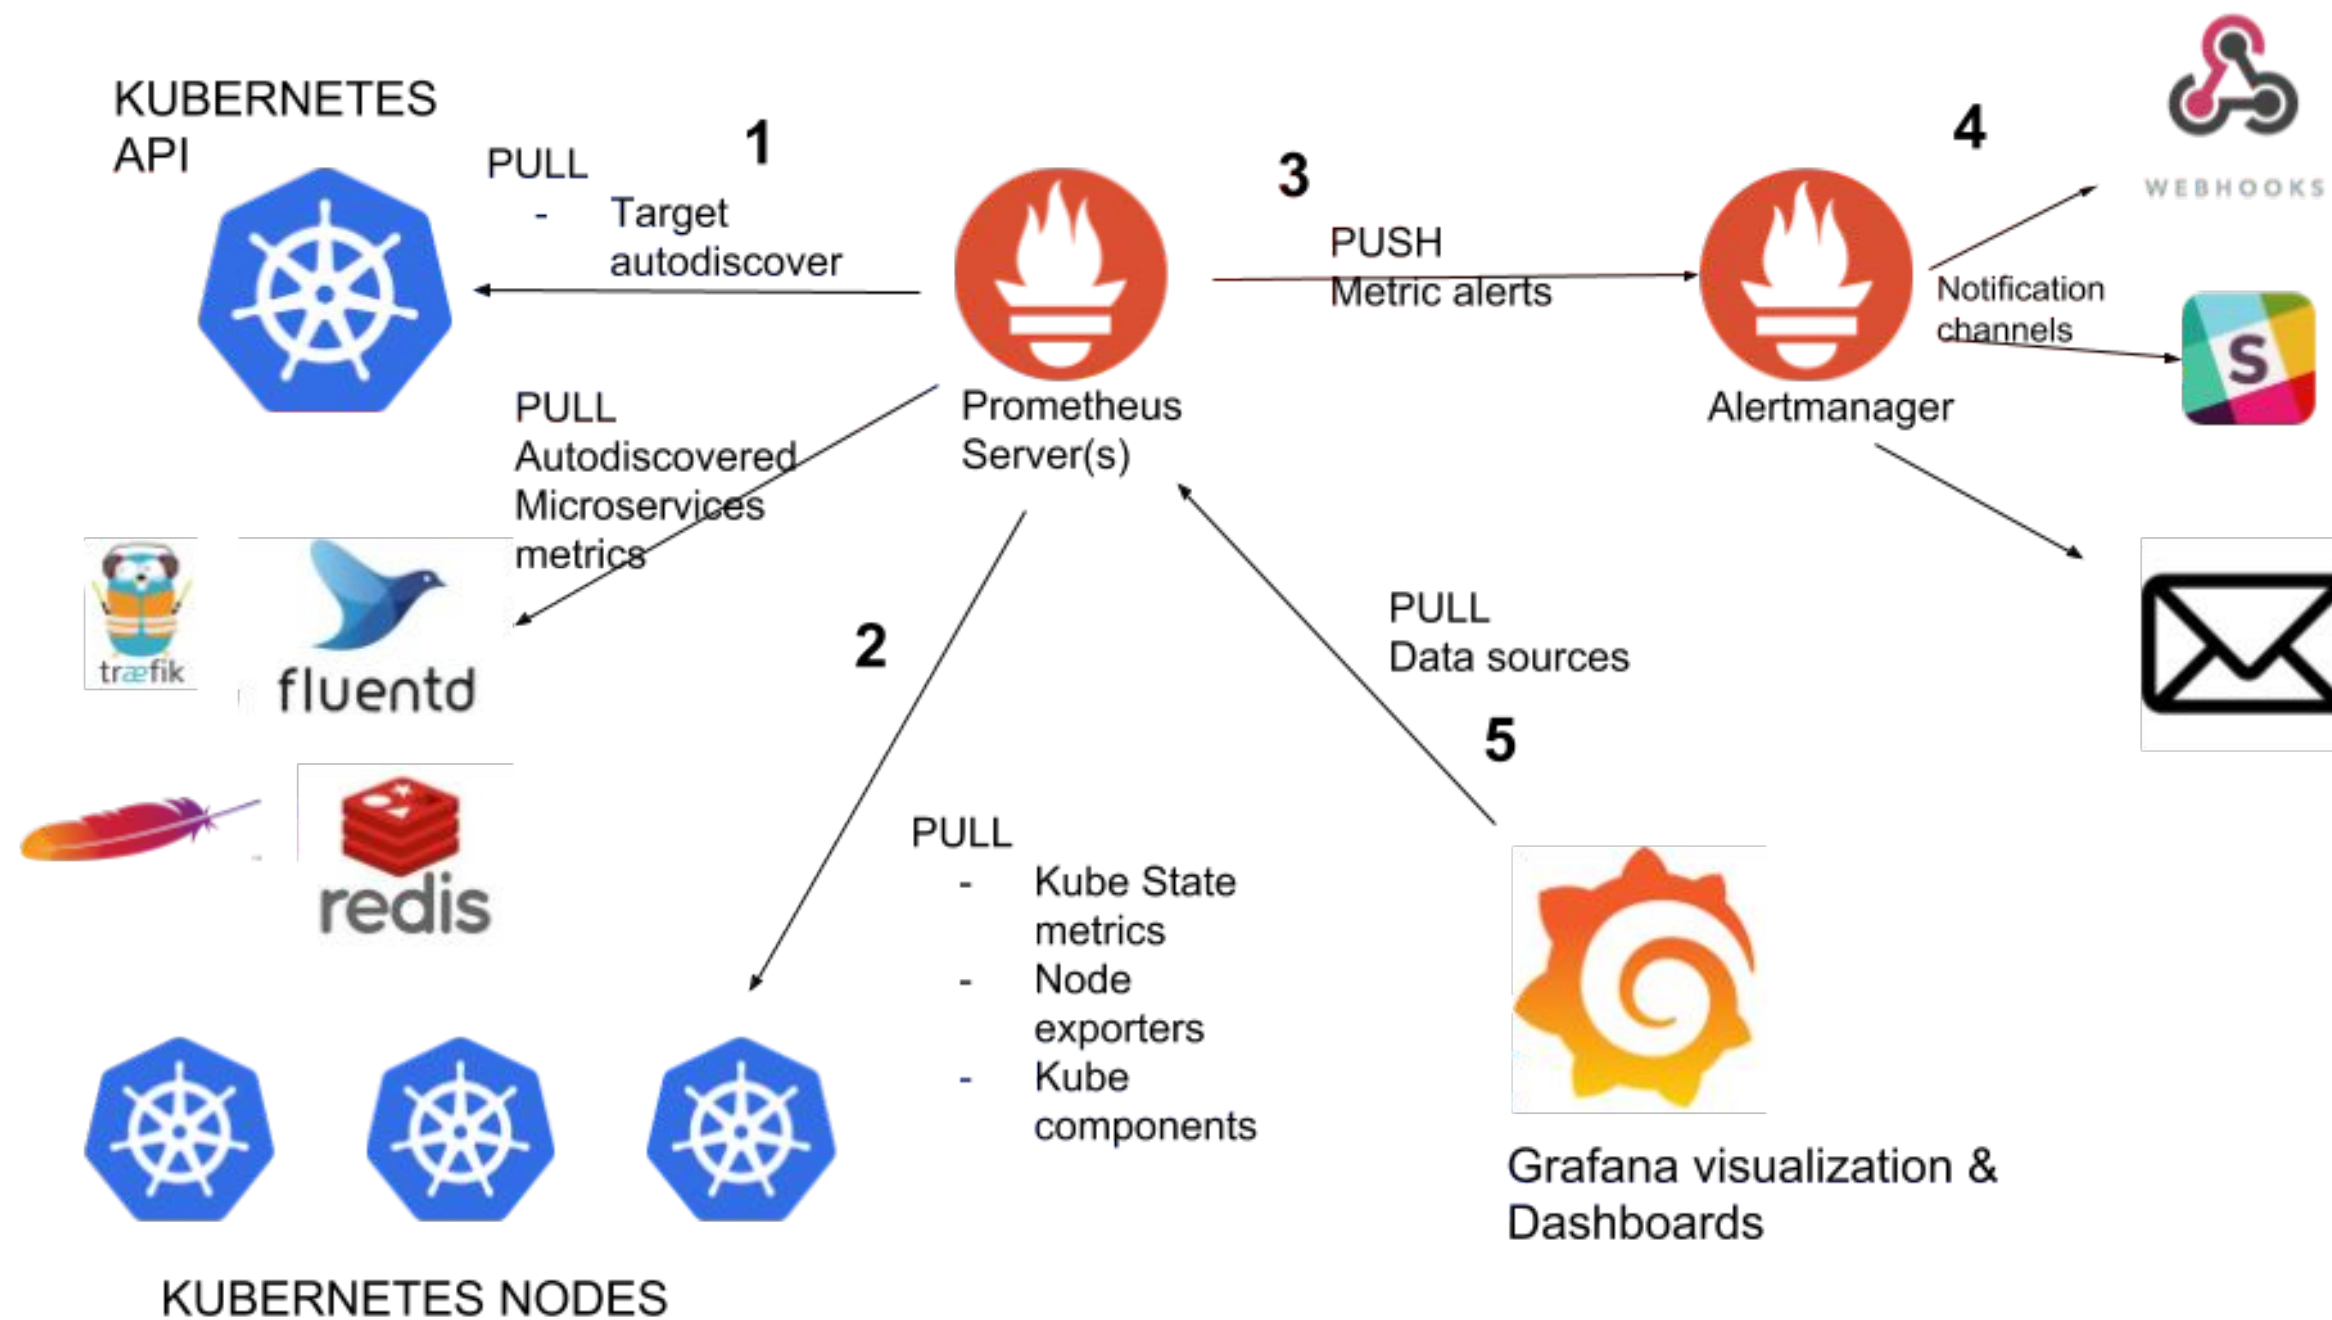
Prometheus:

- It developed as SoundCloud by ex-Googleers
- Prometheus is a close cousin of Kubernetes
- A multi-dimensional data model with time series data identified by metric name and key/value pairs
- Alerting and graphing are unified, using the same language
- Time series collection happens via a pull model over HTTP
- Targets are discovered via service discovery or static configuration
- Provides multiple exporters to send AWS EC2, Kafka, Redis, RMQ,... metrics
- You can export specific metrics from Prometheus to Kubernetes and use these metrics in HPA

Monitoring



Monitoring



Monitoring

Metrics:

- CPU (system, user, nice, iowait, idle, irq, softirq, guest)
- Memory (Apps, Buffers, Cached, Free, SwapCached, Active, Inactive,...)
- Load
- Disk Space Used in percent
- Disk Utilization per Device
- Disk IOPS per device (read, write)
- Disk Throughput per Device (read, write)
- Context Switches
- Network Traffic (In, Out)
- UDP Stats
- Contrack

Monitoring

demo_api_request_duration_seconds_count

Execute

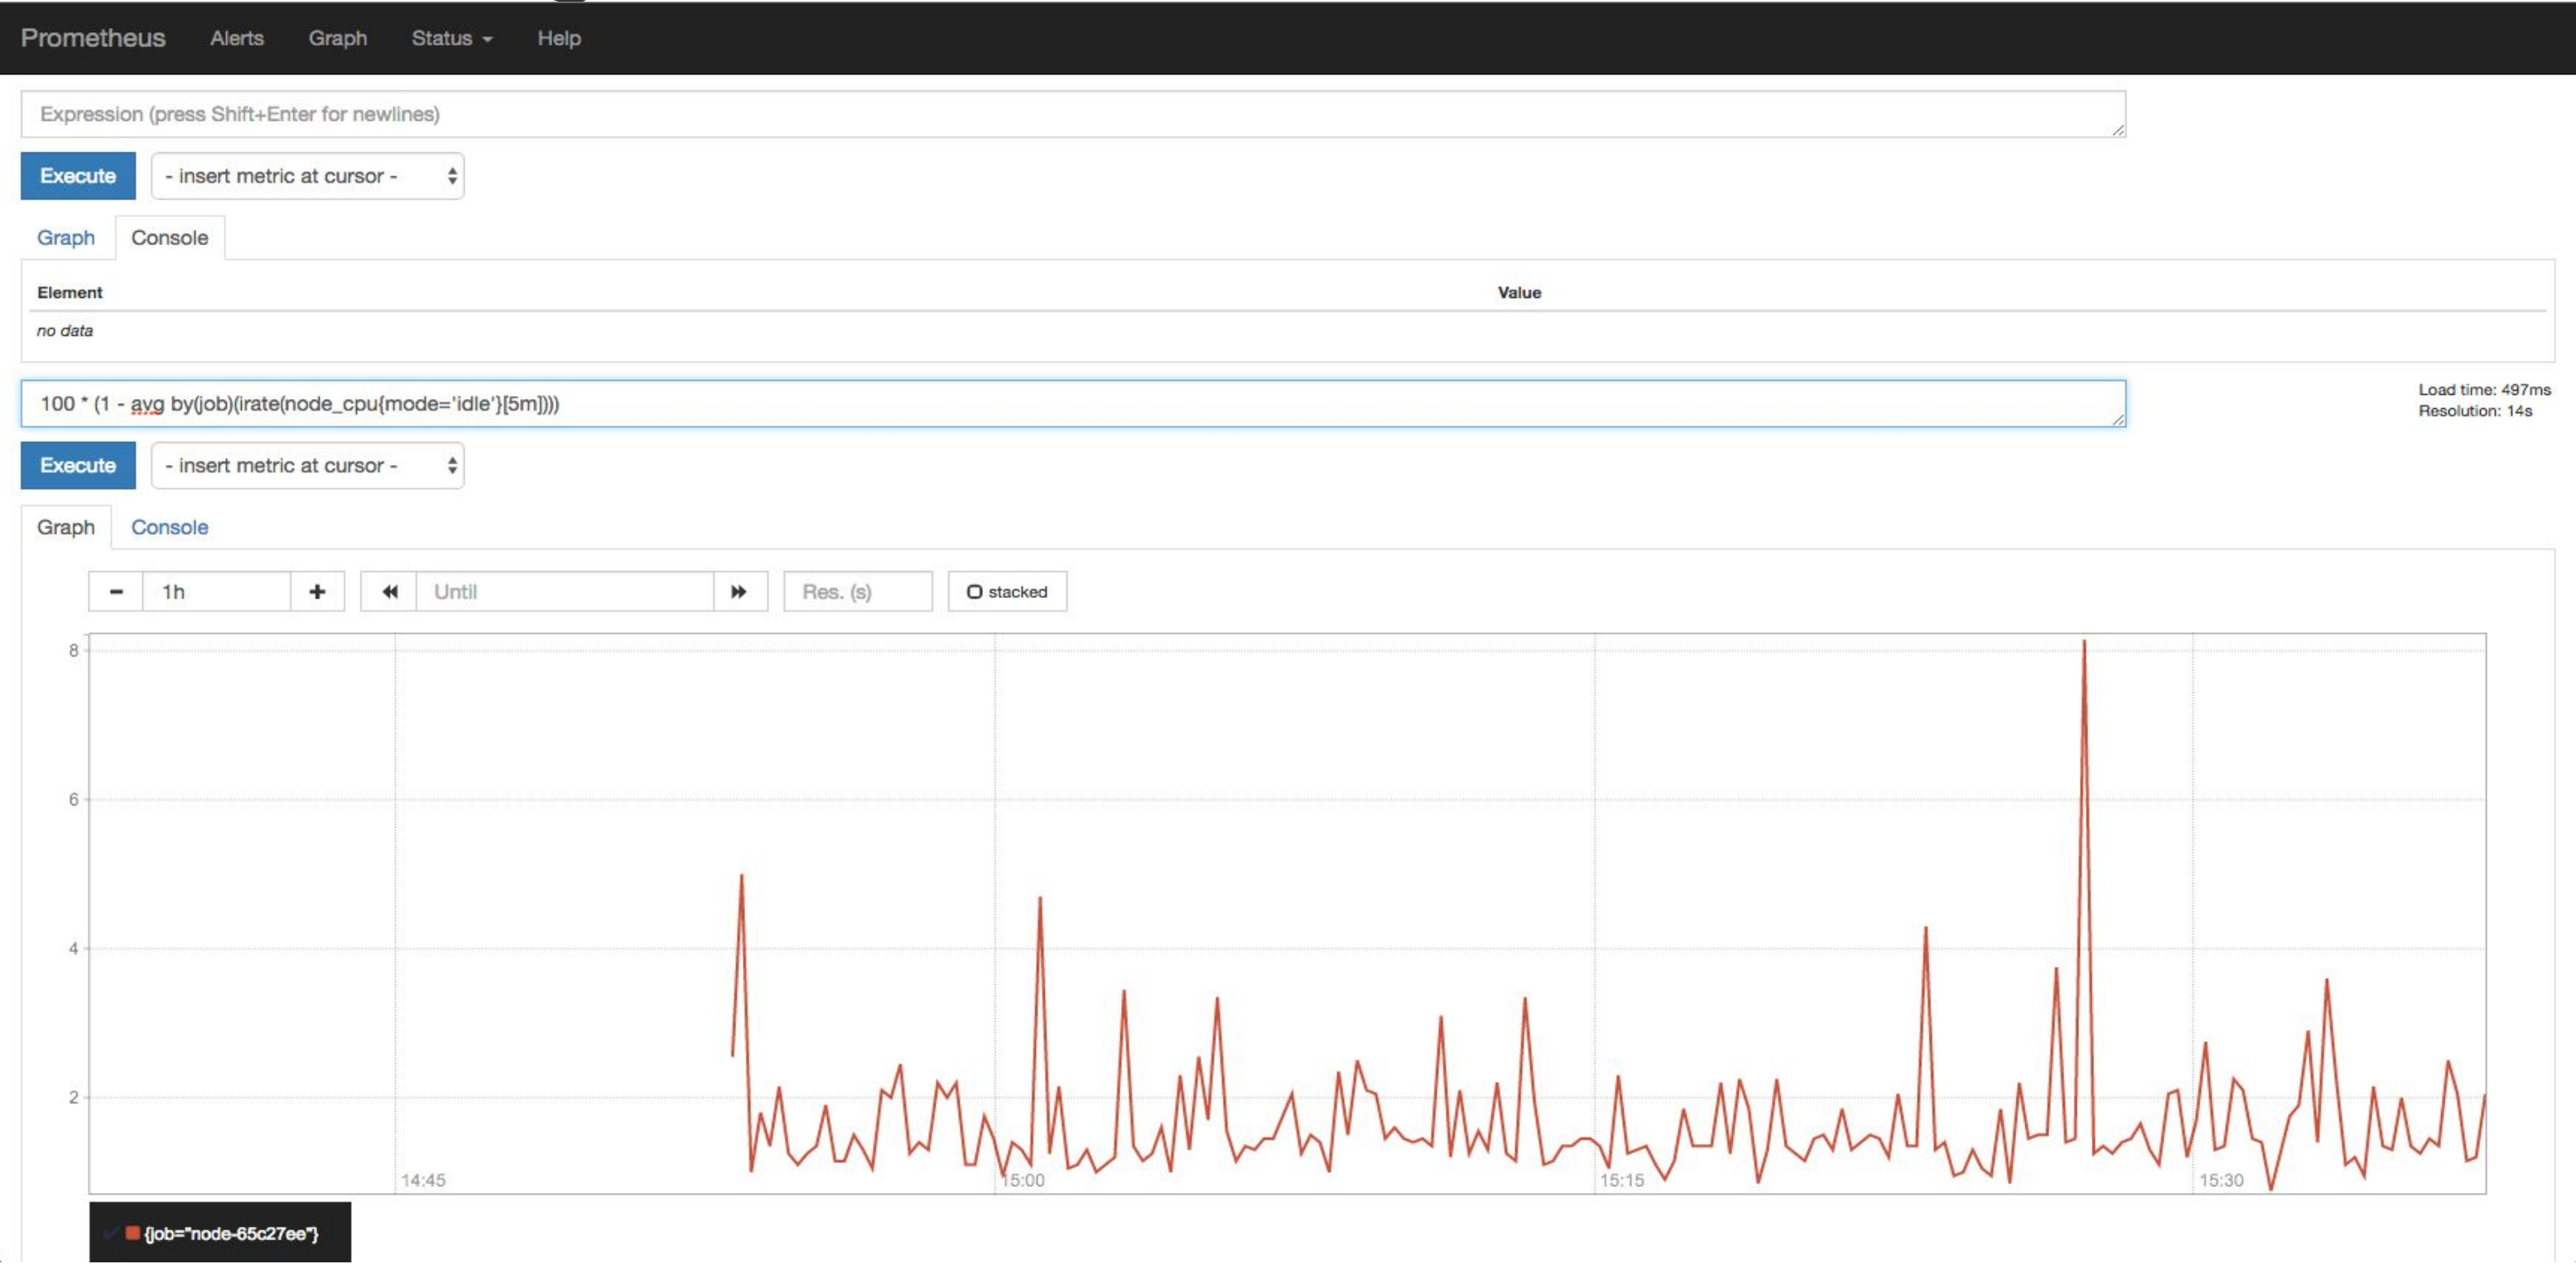
- insert metric at cursor - ▼

Graph

Console

Element	Value
demo_api_request_duration_seconds_count{instance="localhost:8081",job="demo",method="GET",path="/api/nonexistent",status="404"}	216
demo_api_request_duration_seconds_count{instance="localhost:8082",job="demo",method="GET",path="/api/foo",status="500"}	56
demo_api_request_duration_seconds_count{instance="localhost:8080",job="demo",method="POST",path="/api/foo",status="500"}	37
demo_api_request_duration_seconds_count{instance="localhost:8080",job="demo",method="POST",path="/api/foo",status="200"}	632
demo_api_request_duration_seconds_count{instance="localhost:8080",job="demo",method="GET",path="/api/foo",status="200"}	6820
demo_api_request_duration_seconds_count{instance="localhost:8081",job="demo",method="POST",path="/api/bar",status="500"}	19
demo_api_request_duration_seconds_count{instance="localhost:8081",job="demo",method="GET",path="/api/bar",status="200"}	3671

Monitoring



Monitoring

Recap:

- Monitoring is essential to run, understand and operate services.
- Prometheus
 - Client instrumentation
 - Scrape configuration
 - Querying
 - Dashboards
 - Alert rules
- Important Metrics
 - Four golden signals: Latency, Traffic, Error, Saturation
- Best Practices

Logging

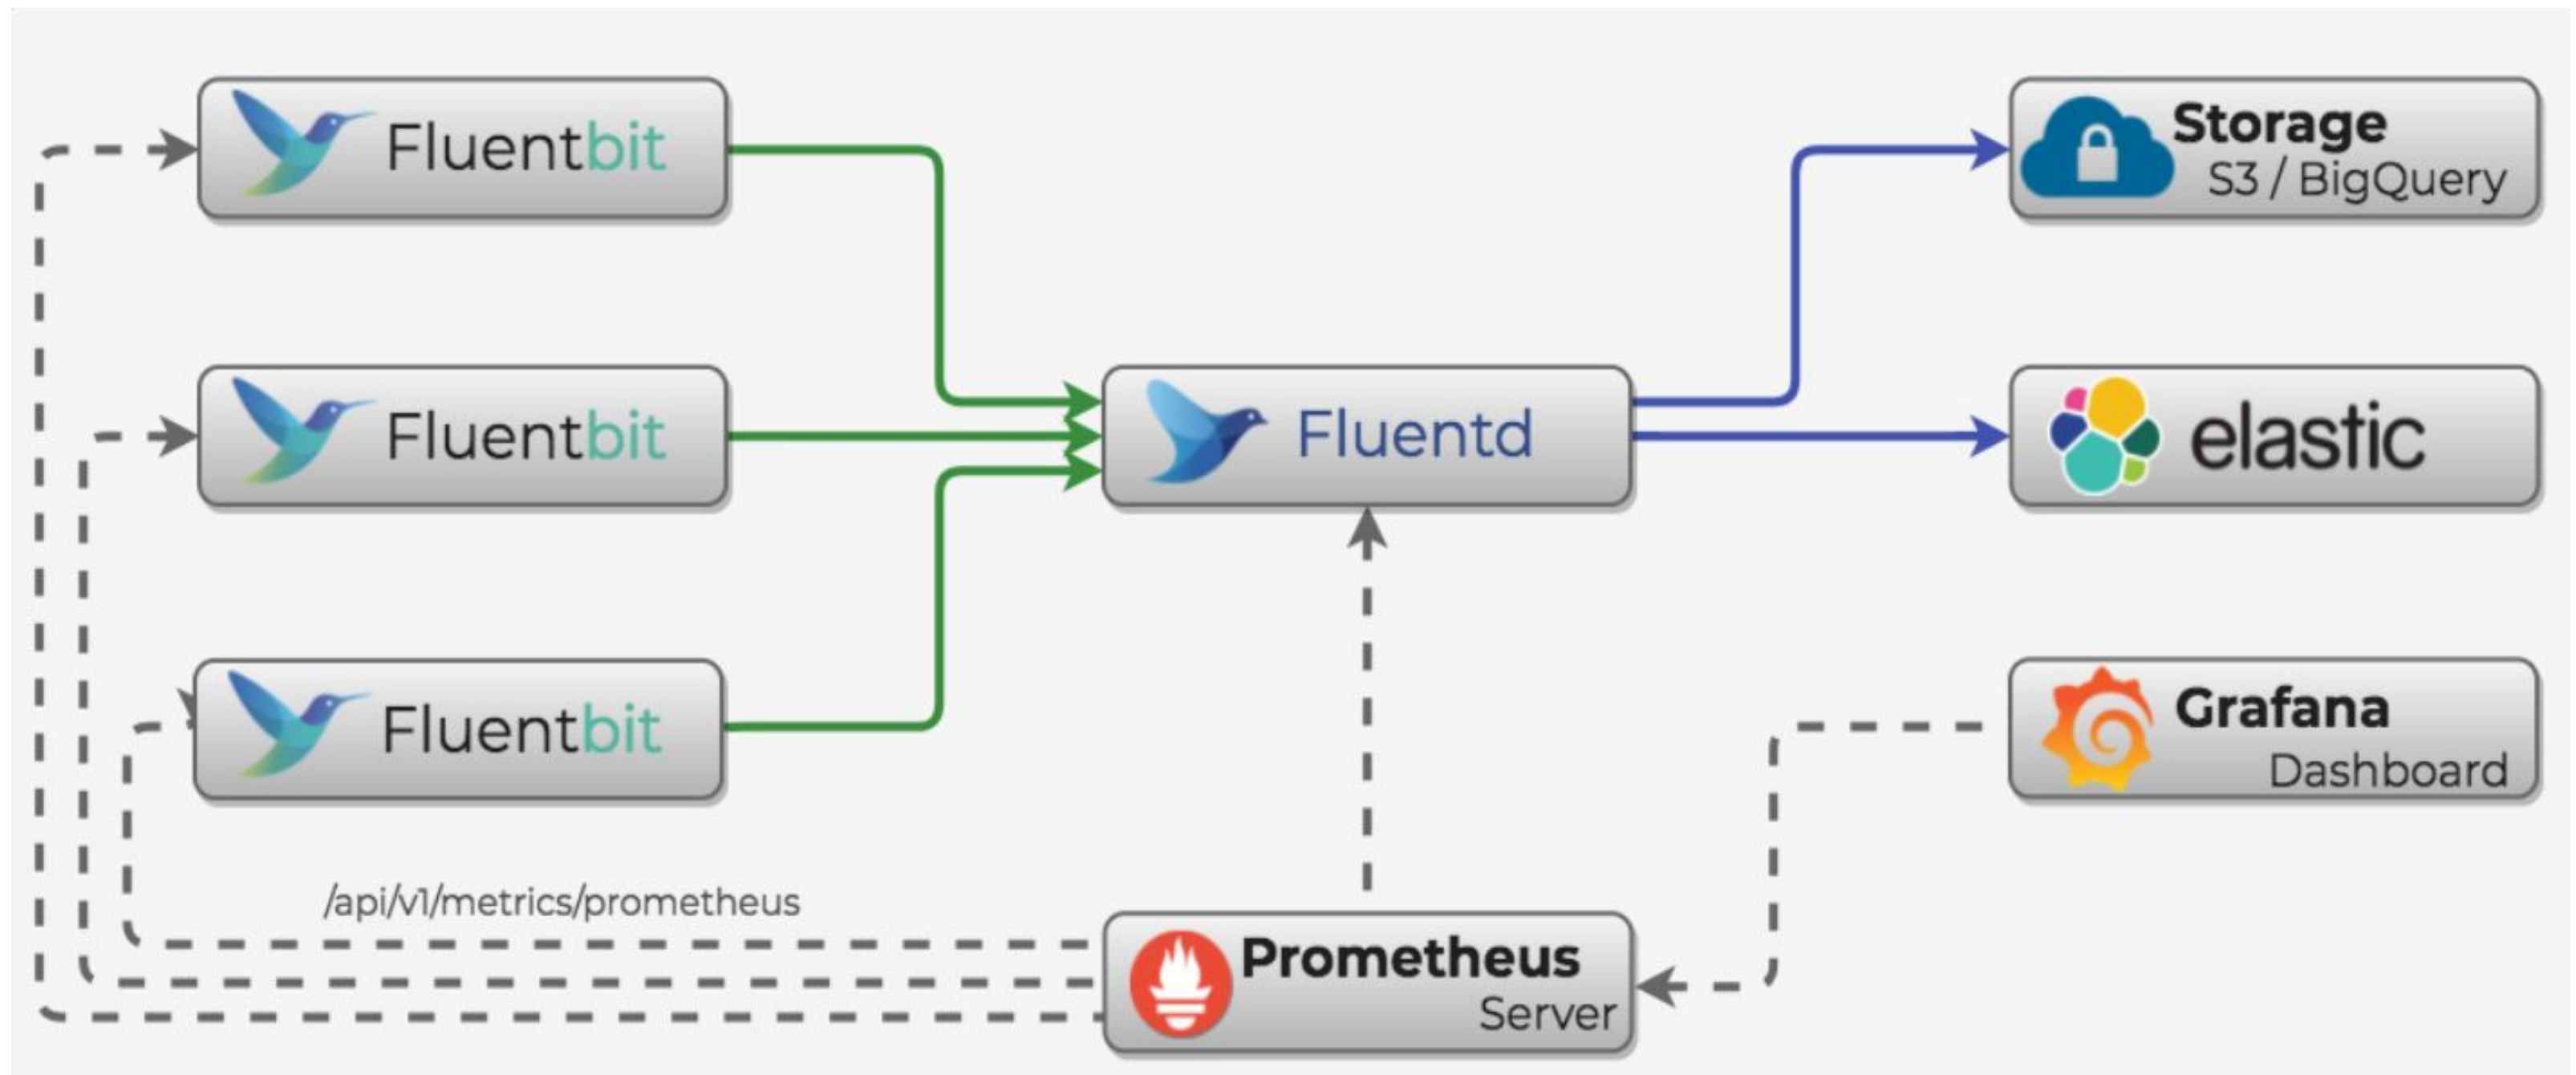
Your pods logs exist on the node only as long as the pod itself exists on the Kubernetes node; this means, when there's an eviction event — for example, rescheduling pods from a tainted node-, and the pods are drained, the logs are taken with it. Persisting this data, however, can be useful, and doing this requires a little bit of iteration on the logging architecture. Persistence aside, this also allows you to leverage more advanced tooling around your log data once aggregated.

How your cluster ships logs to a logging backend will vary in small ways, but typically will come down to the following process; a DaemonSet (meaning a single pod that will run on as many nodes that exist in the cluster; if you have 4 nodes, you have 4 instances, and if you remove a node, that 4th instance is not rescheduled) is used to run a logging agent, and that agent will ship logs to your logging backend.

Logging

Unfortunately, Kubernetes doesn't provide many configuration options that pertain to logging. Docker offers multiple logging [drivers](#), but we can't configure them via Kubernetes. Thankfully, there are already several open source solutions. When it comes to logging, our favorite tools are **Fluentd**, **Fluent-bit** and **Filebeat**.

Logging



Logging: Fluentd

Fluentd is an open source log collector, processor, and aggregator that was created back in 2011 by the folks at Treasure Data. Written in Ruby, Fluentd was created to act as a unified logging layer — a one-stop component that can aggregate data from multiple sources, unify the differently formatted data into JSON objects and route it to different output destinations.

Design wise — performance, scalability, and reliability are some of Fluentd's outstanding features. A vanilla Fluentd deployment will run on ~40MB of memory and is capable of processing above 10,000 events per second. Adding new inputs or outputs is relatively simple and has little effect on performance. Fluentd uses disk or memory for buffering and queuing to handle transmission failures or data overload and supports multiple configuration options to ensure a more resilient data pipeline.

Logging: Logstash

Logstash is not the oldest shipper of this list (that would be syslog-ng, ironically the only one with “new” in its name), but it’s certainly **the best known**. That’s because **it has lots of plugins**: inputs, codecs, filters and outputs. Basically, you can take pretty much any kind of data, enrich it as you wish, then push it to lots of destinations.

Logstash Advantages:

Logstash’s main strongpoint is **flexibility, due to the number of plugins**.

Logstash Disadvantages:

Logstash’s biggest con or “Achille’s heel” has always been **performance and resource consumption** (the default heap size is 1GB).

Logging: Fluentd

Fluentd has been around for some time now and has developed a rich ecosystem consisting of more than 700 different plugins that extend its functionality. Fluentd is the de-facto standard log aggregator used for logging in Kubernetes and as mentioned above, is one of the widely used Docker images.

Logging: Fluent Bit

Fluent Bit is an open source log collector and processor also created by the folks at Treasure Data in 2015. Written in C, Fluent Bit was created with a specific use case in mind — highly distributed environments where limited capacity and reduced overhead (memory and CPU) are a huge consideration.

To serve this purpose, Fluent Bit was designed for high performance and comes with a super light footprint, running on ~450KB only. An abstracted I/O handler allows asynchronous and event-driven read/write operations. For resiliency and reliability, various configuration options are available for defining retries and the buffer limit.

Logging: Fluent Bit

Fluent Bit is also extensible, but has a smaller eco-system compared to Fluentd. Inputs include syslog, tcp, systemd/journald but also CPU, memory, and disk. Outputs include Elasticsearch, InfluxDB, file and http. For Kubernetes deployments, a dedicated filter plugin will add metadata to log data, such as the pod's name and namespace, and the containers name/ID.

Logging: Filebeat

Filebeat is a lightweight logshipper for logstash. It can be installed as agents on your servers to collect operational data. Since it is lightweight it does not consume system resources compared to Logstash. It has a number of modules that can help in parsing, collecting and visualization of logs. Moreover, it is container ready so in case if you are looking to build a container based solution you can run Filebeat inside a container in the same Virtual Machine

Logging: Compare

	Fluentd	Fluent Bit	Filebeat
Scope	Containers/Servers	Containers/Servers	Containers/Servers
Language	C&Ruby	C	GO
memory on start	~40MB	~450KB	~50-200MB
Performance	High Performance	High Performance	High Performance
Dependencies	Built as a Ruby Gem, It requires a certain number of gems	Zero dependencies, unless some special plugin requires them	Zero dependencies
Plugins	More than 650 plugins available	Around 35 plugins available	17 internal modules
License	Apache Licence 2.0	Apache Licence 2.0	Apache Licence 2.0

Logging: Conclusion

Logstash, fluentd, Fluent Bit, Filebeat have their own features as follows:

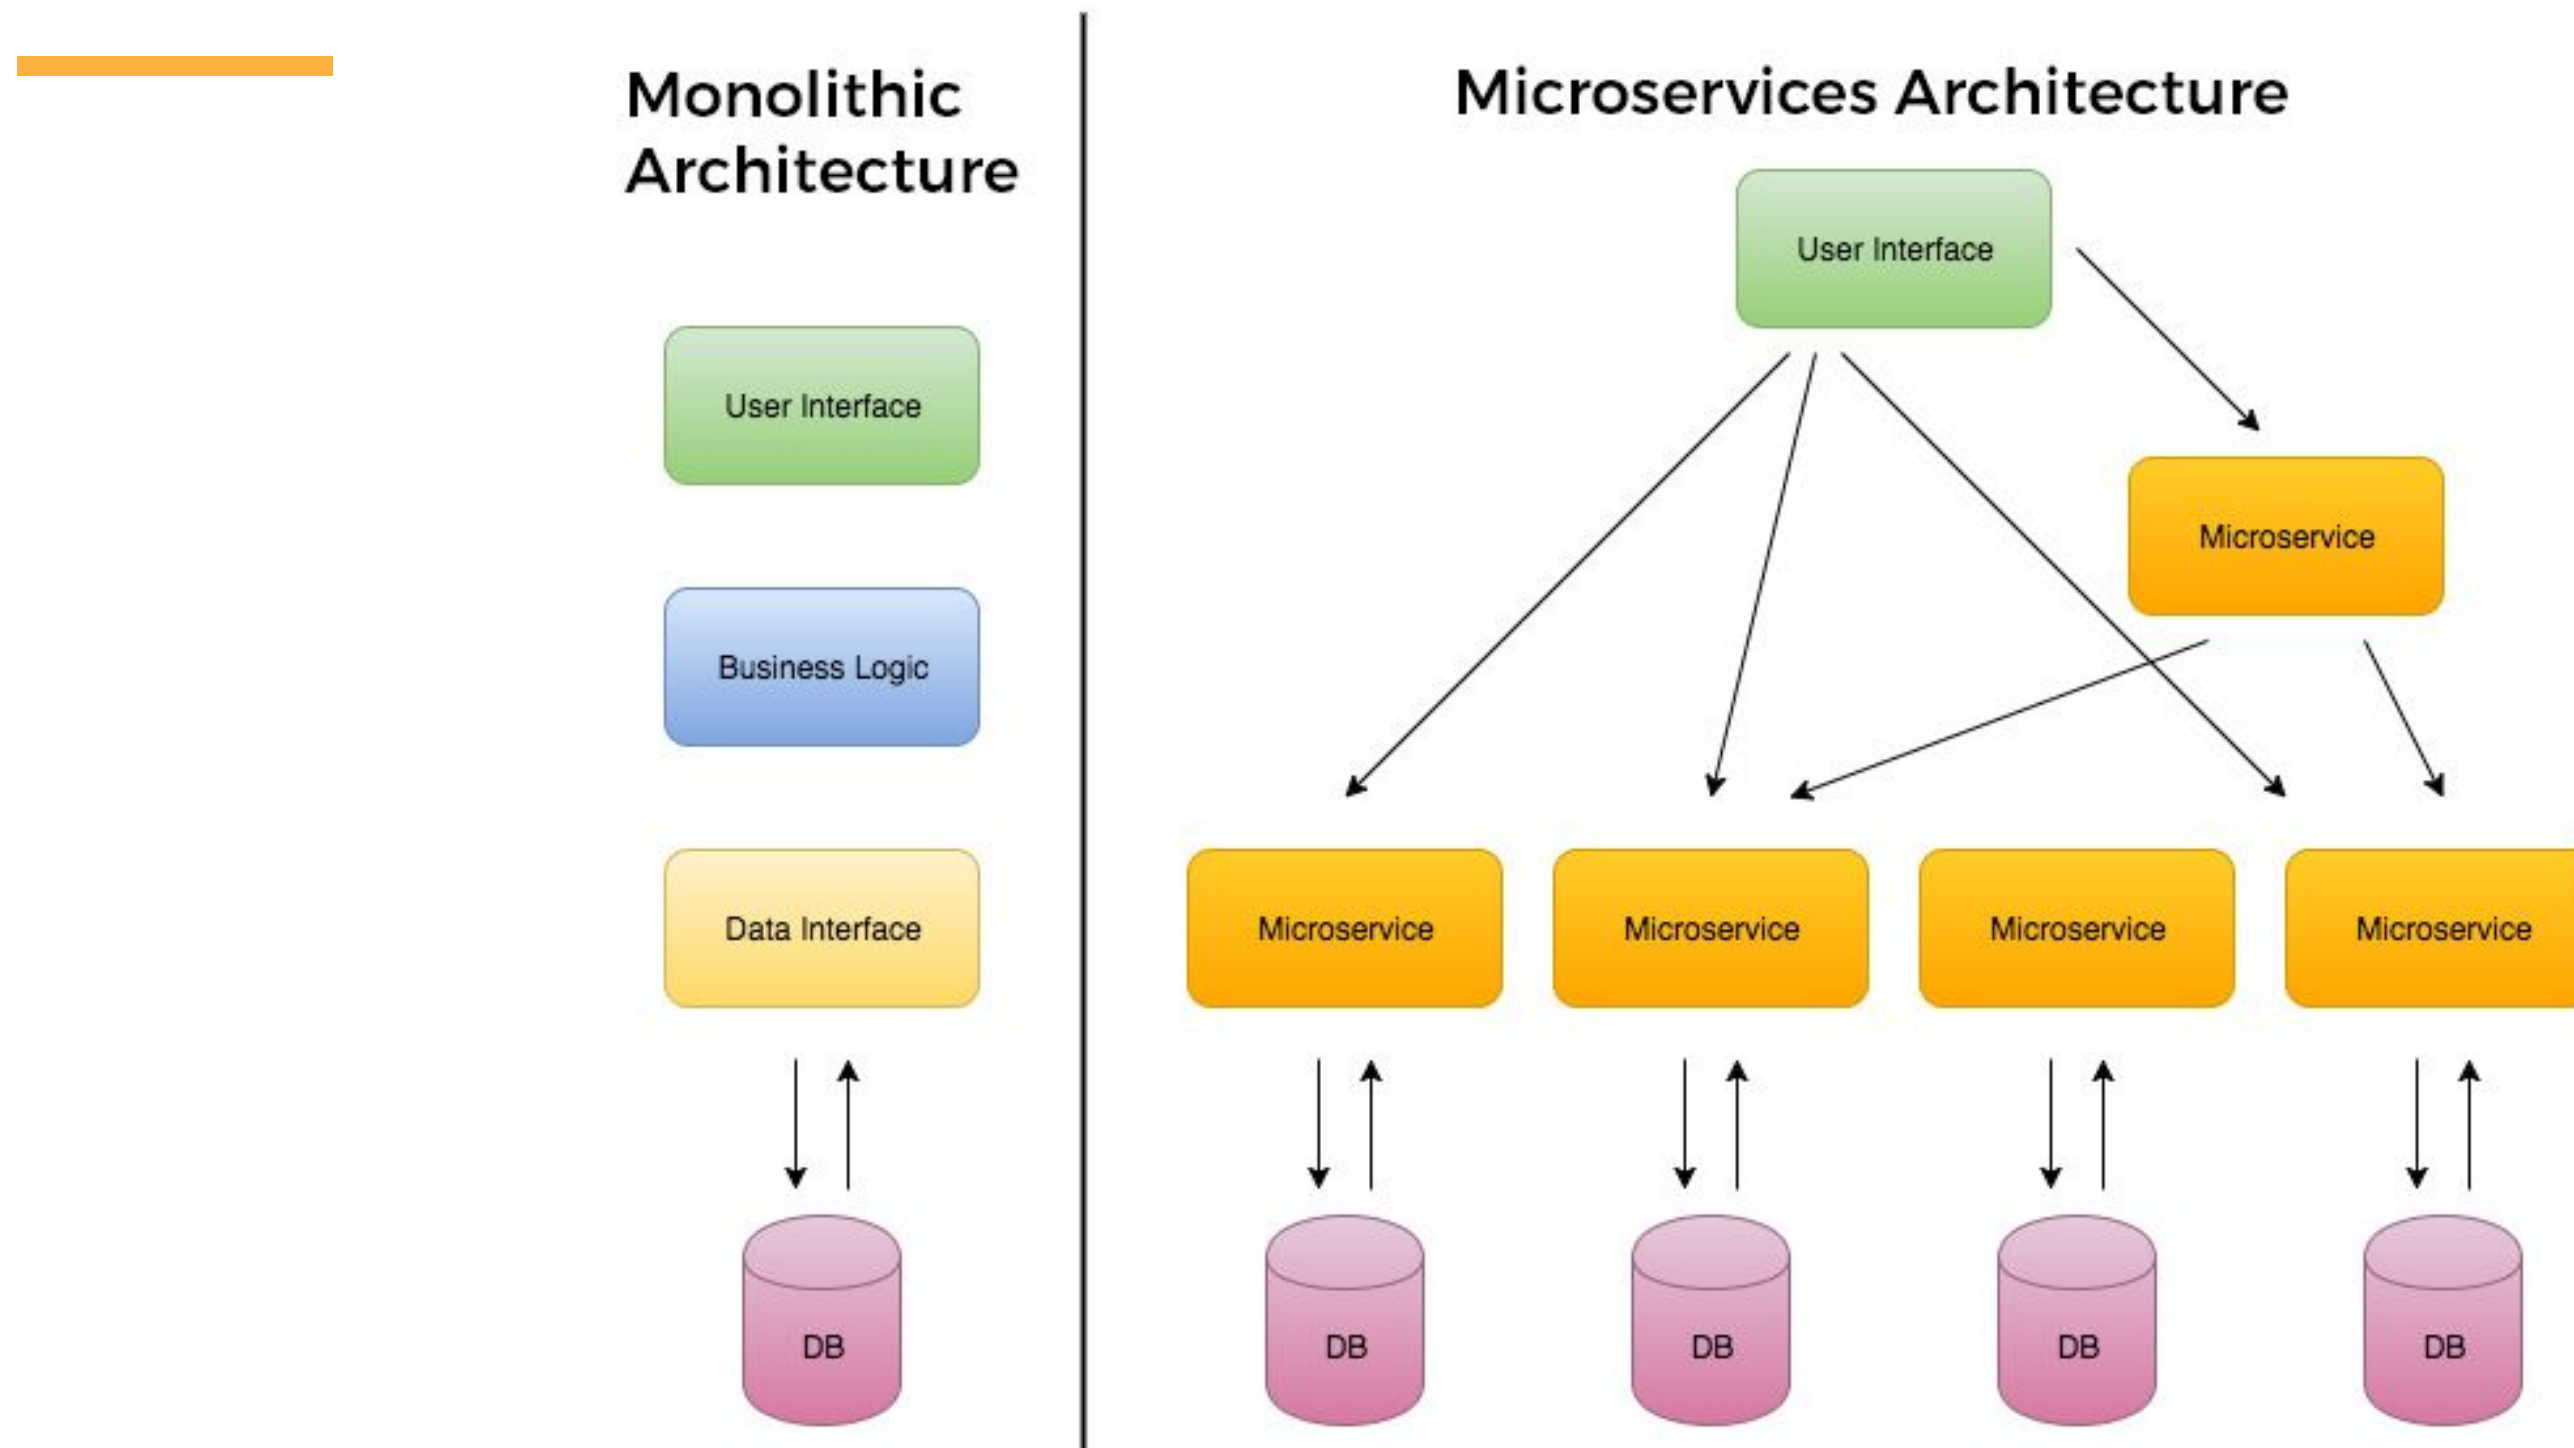
- . Logstash supports all the mainstream log types, diverse plug-ins, and flexible customization, but has relatively low performance and is prone to high memory usage because of JVM.
- . Fluentd supports all the mainstream log types and many plug-ins, and delivers good performance.
- . Fluent bit doesn't has small number of plugins but is very small and has good performance
- . Filebeat - small but has many issues about memory leaks!!!

Distributed tracing systems

Monolithic service architectures for large backend applications are becoming increasingly rare. The monoliths are being replaced with distributed system architectures, where the backend application is spread out (distributed) in an ecosystem of small and narrowly-focused services. These distributed services communicate with each other over the network to process requests.

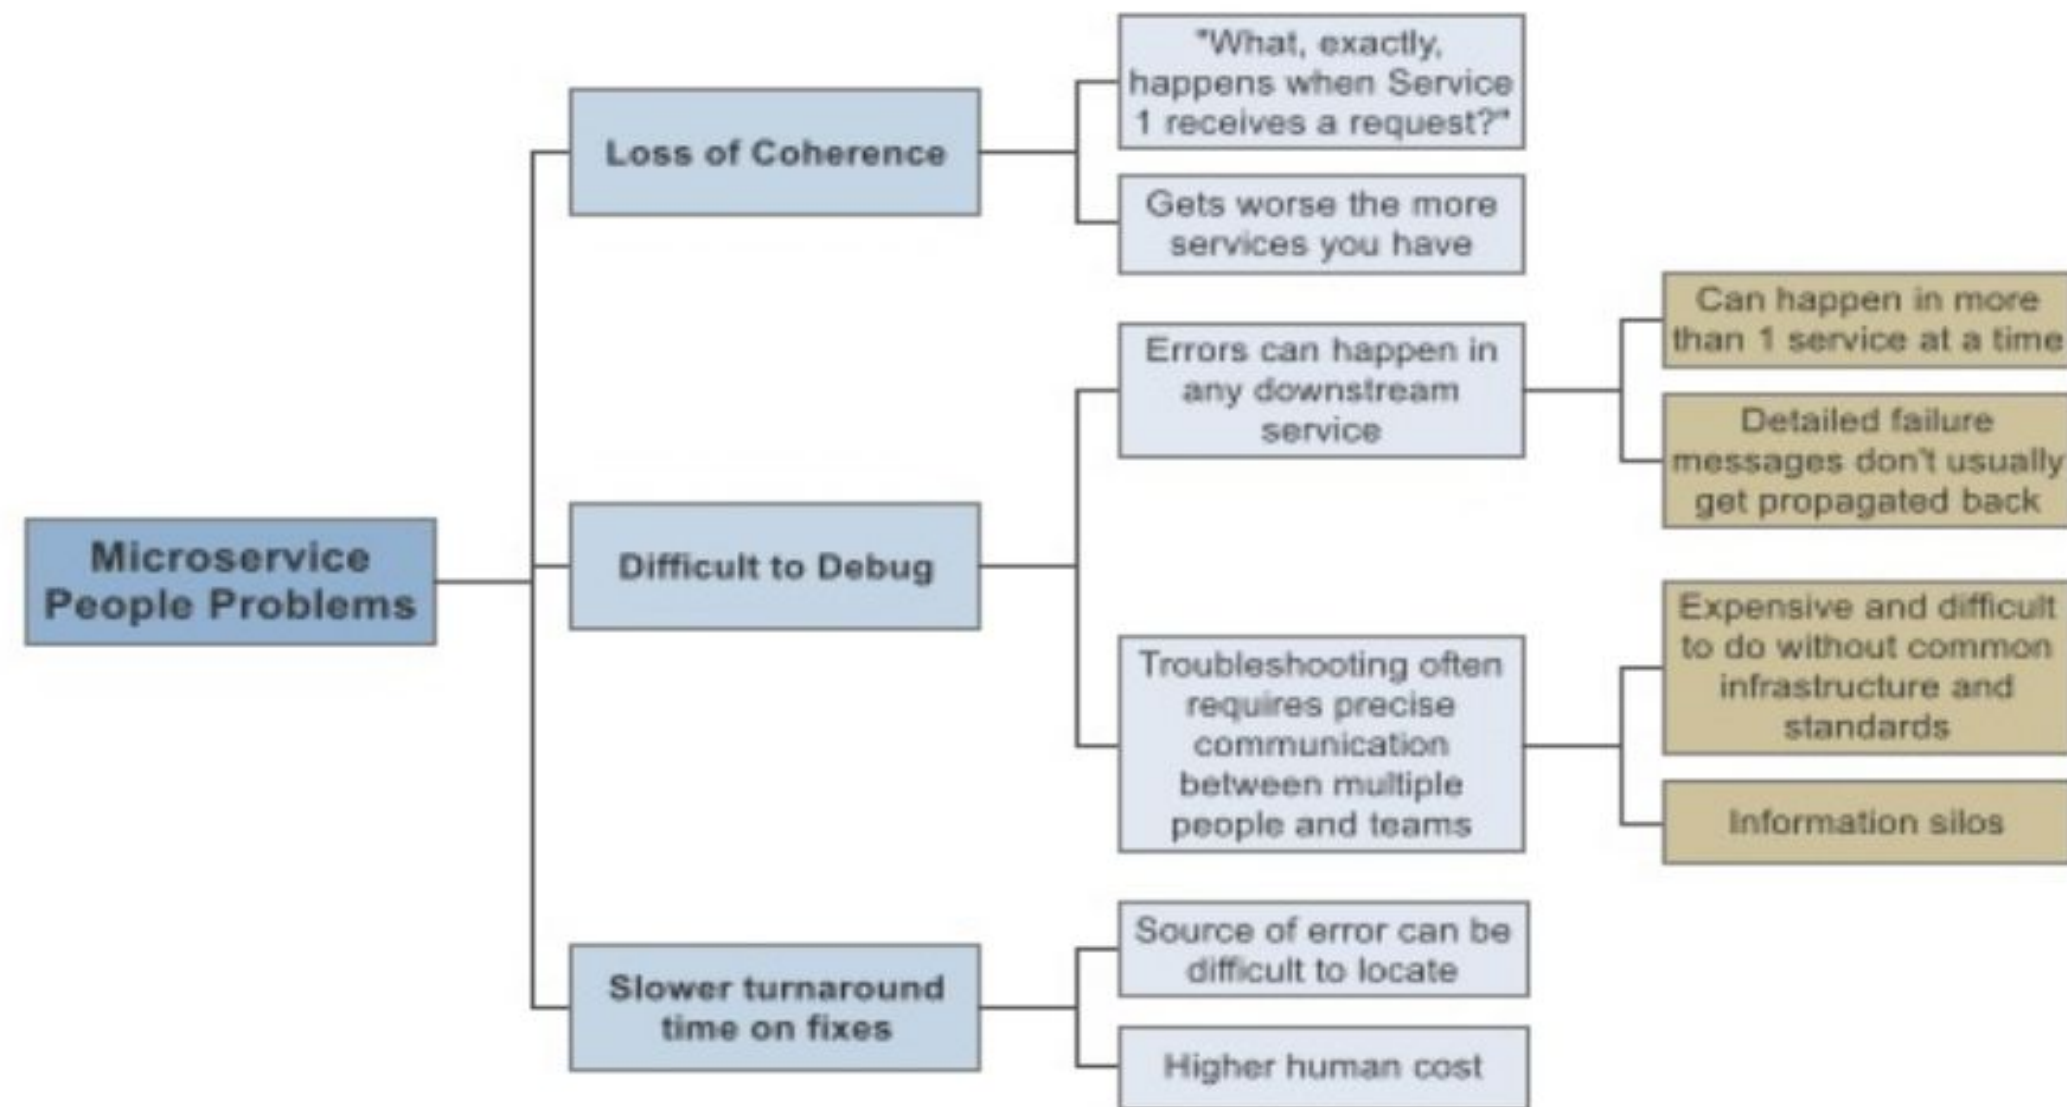
There are many benefits to distributed system architectures, but we're not trying to sell you on them in this blog post. A quick [microservice vs. monolith Google search](#) can do that. Instead, we're going to focus on the difficulty of tracking and analyzing requests in a distributed environment; pinpointing problems can be very frustrating when requests might touch dozens of services during processing.

Distributed tracing systems



Distributed tracing systems

— Microservice People Problems



Distributed tracing systems

Open source distributed tracing tools:

1. Zipkin - Twitter. Now - CNCF
2. Jaeger - Uber. Now - CNCF
3. Opencensus - Google
4. Opentracing - CNCF



Distributed tracing systems

Zipkin and Jaeger are two popular choices for request tracing. Zipkin was originally inspired by Dapper and developed by Twitter. It's now maintained by a dedicated community.

Jaeger was originally built and open sourced by Uber.

Jaeger is a Cloud Native Computing Foundation project. The overall architecture is similar. Instrumented systems send events/traces to the trace collector.

Distributed tracing systems

OpenTracing is a spec that grew out of Zipkin to provide cross-platform compatibility. It offers a vendor-neutral API for adding tracing to applications and delivering that data into distributed tracing systems. A library written for the OpenTracing spec can be used with any system that is OpenTracing-compliant.

OpenCensus is a vendor-agnostic (unified framework) single distribution of libraries to provide **metrics** collection and **tracing** for your services. *OpenCensus* originates from *Google*.

Distributed tracing systems: Zipkin vs Jaeger

	JAEGER	ZIPKIN
OpenTracing compatibility	Yes	Yes
Language support	Python Go Node Python Java C++ C# Ruby PHP	Go Java JavaScript Ruby Scala C++ C# Python PHP
Storage support	In-memory Cassandra Elasticsearch ScyllaDB (work in progress)	In-memory Cassandra Elasticsearch Mysql

Distributed tracing systems: Zipkin vs Jaeger

	JAEGER	ZIPKIN
Sampling	Dynamic sampling rate (supports rate limiting and probabilistic sampling strategies)	Fixed sampling rate (supports probabilistic sampling strategy)
SPAN transport	UDP HTTP	HTTP Kafka Scribe AMQP
Docker ready	Yes	Yes
Jaeger Agent in Kubernetes as DaemonSet	Yes	No

Distributed tracing systems: SPAN

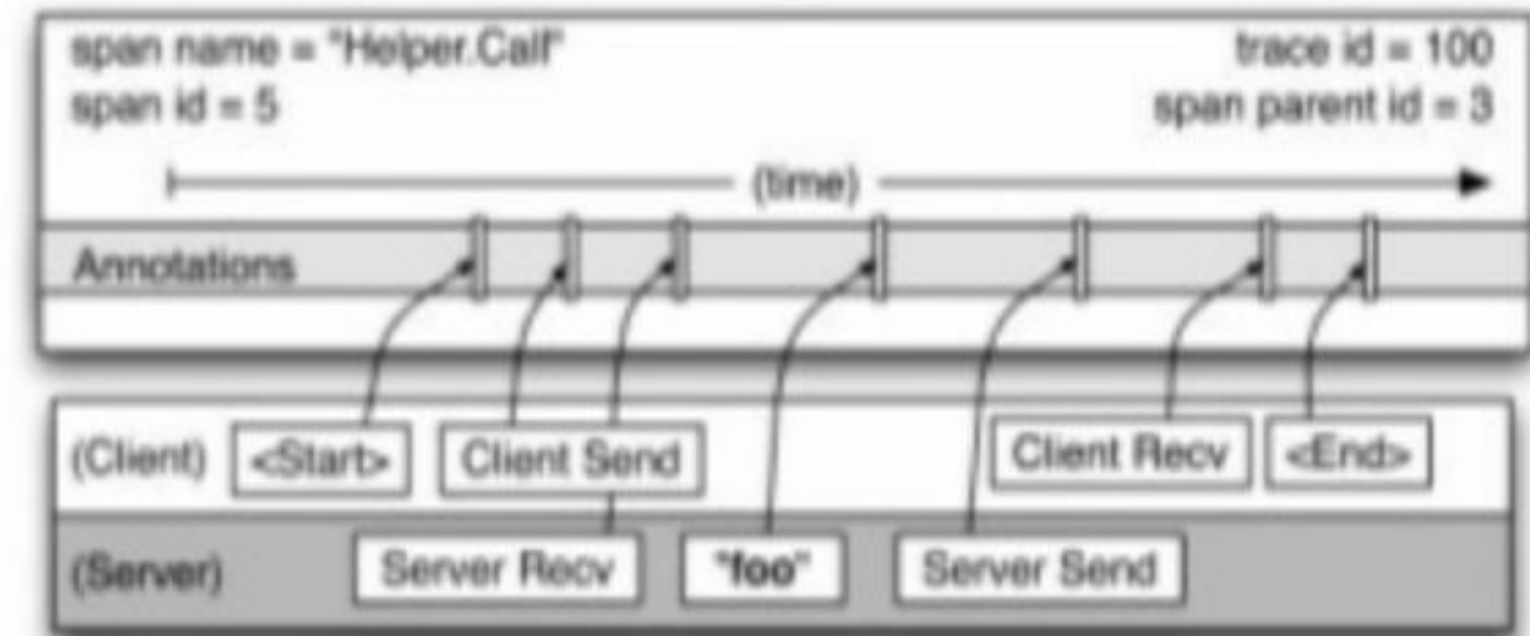
A set of Annotations and BinaryAnnotations that correspond to a particular RPC. Spans contain identifying information such as traceId, spanId, parentId, and RPC name.

Spans are usually small. For example, the serialized form is often measured in KiB or less. When spans grow beyond orders of KiB, other problems occur, such as hitting limits like Kafka message size (1MiB). Even if you can raise message limits, large spans will increase the cost and decrease the usability of the tracing system. For this reason, be conscious to store data that helps explain system behavior, and don't store data that doesn't.

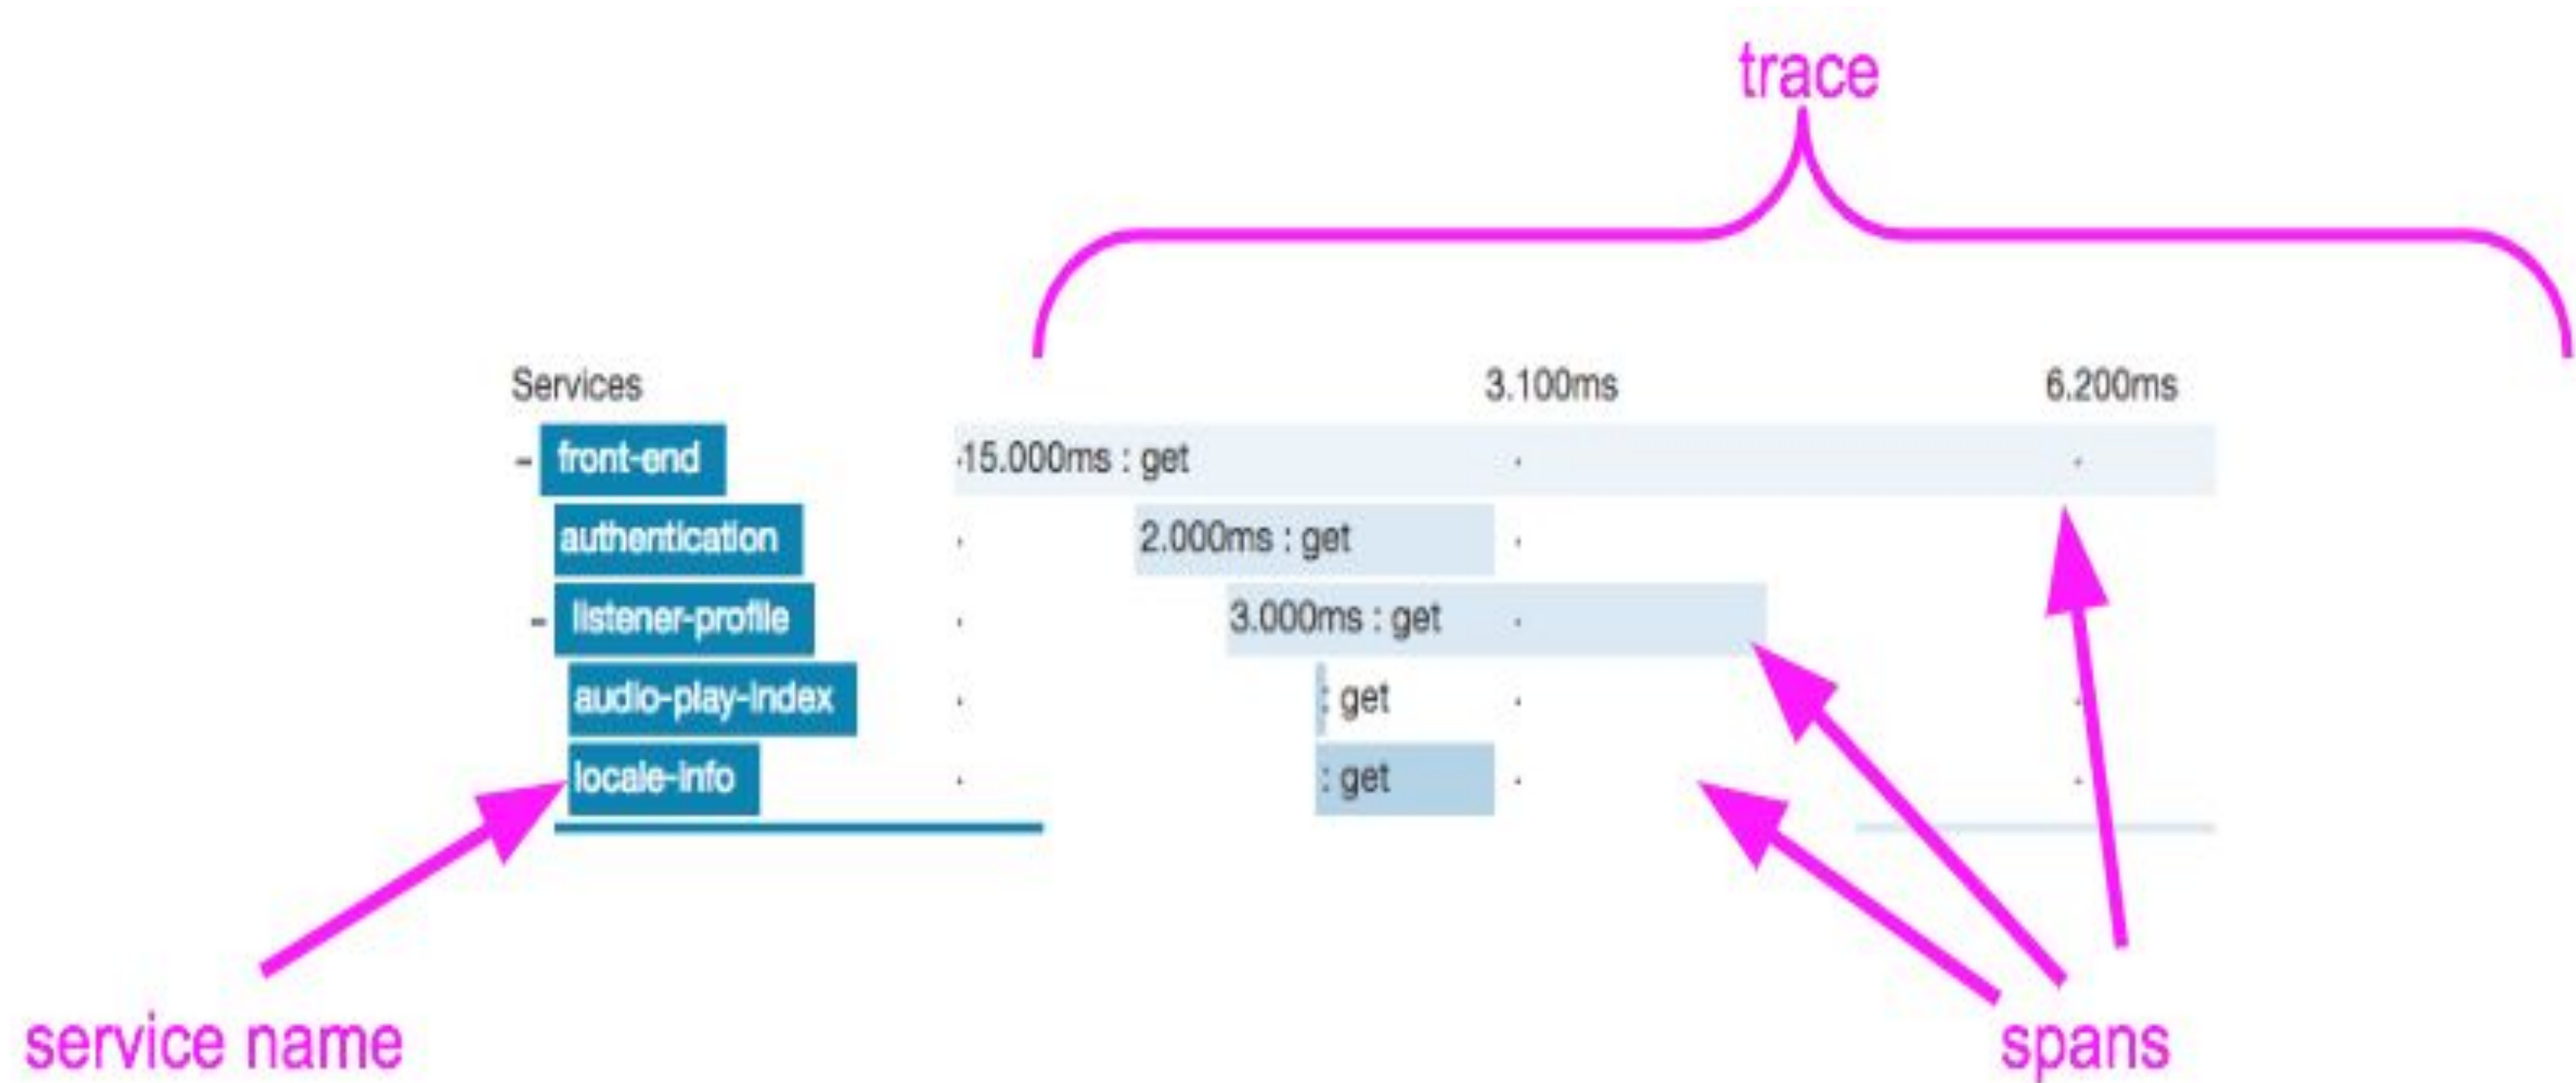
Distributed tracing systems: SPAN

Logged event in a typical span

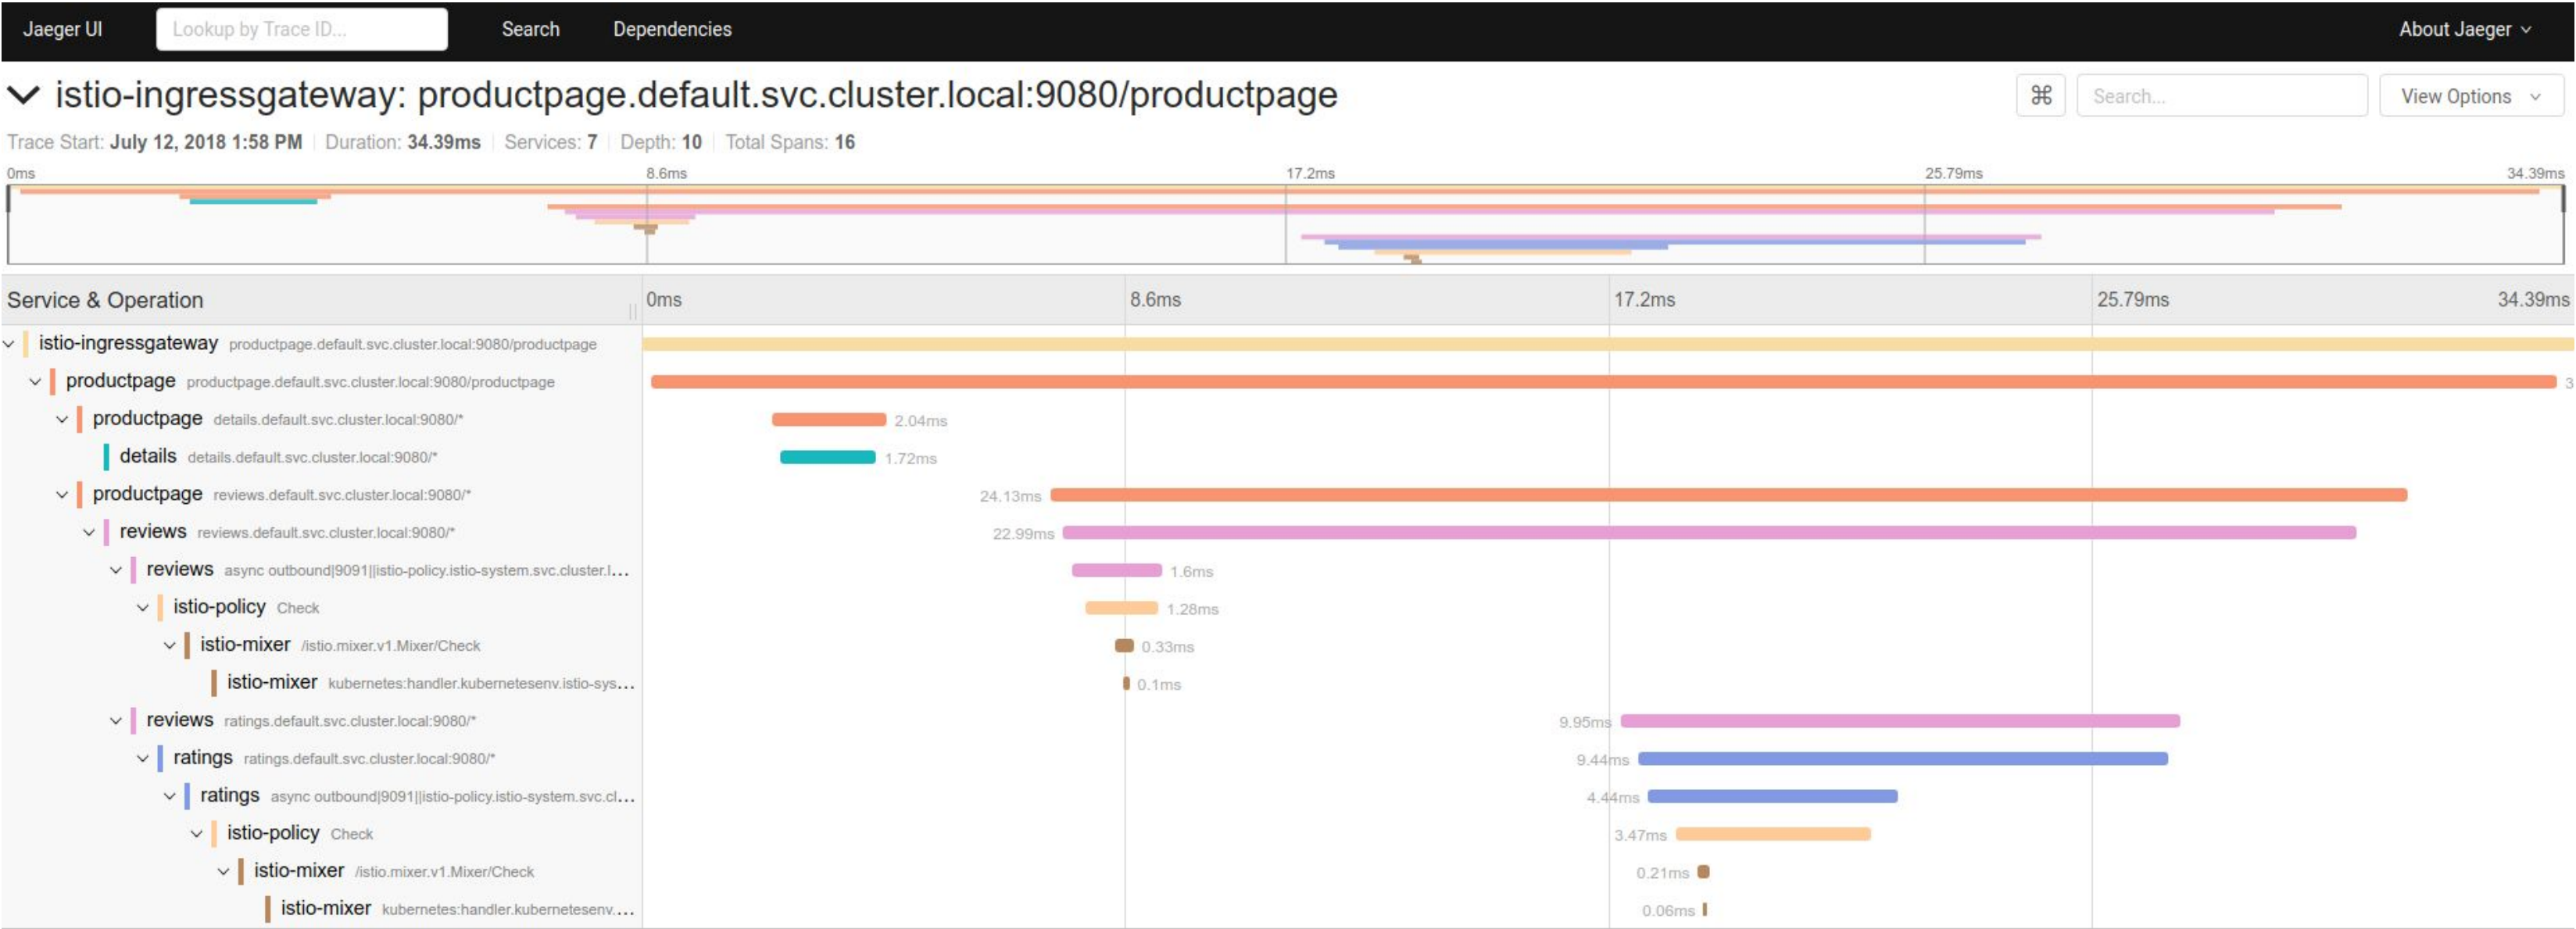
- Span name
- Span start time
- Span end time
- Trace id
- Span id
- Span parent id
- Any timing information recorded by the instrumentation library (RPC, HTTP)
- Additional custom labels ("foo")



Distributed tracing systems: SPAN



Distributed tracing systems




Distributed tracing systems: Opencensus

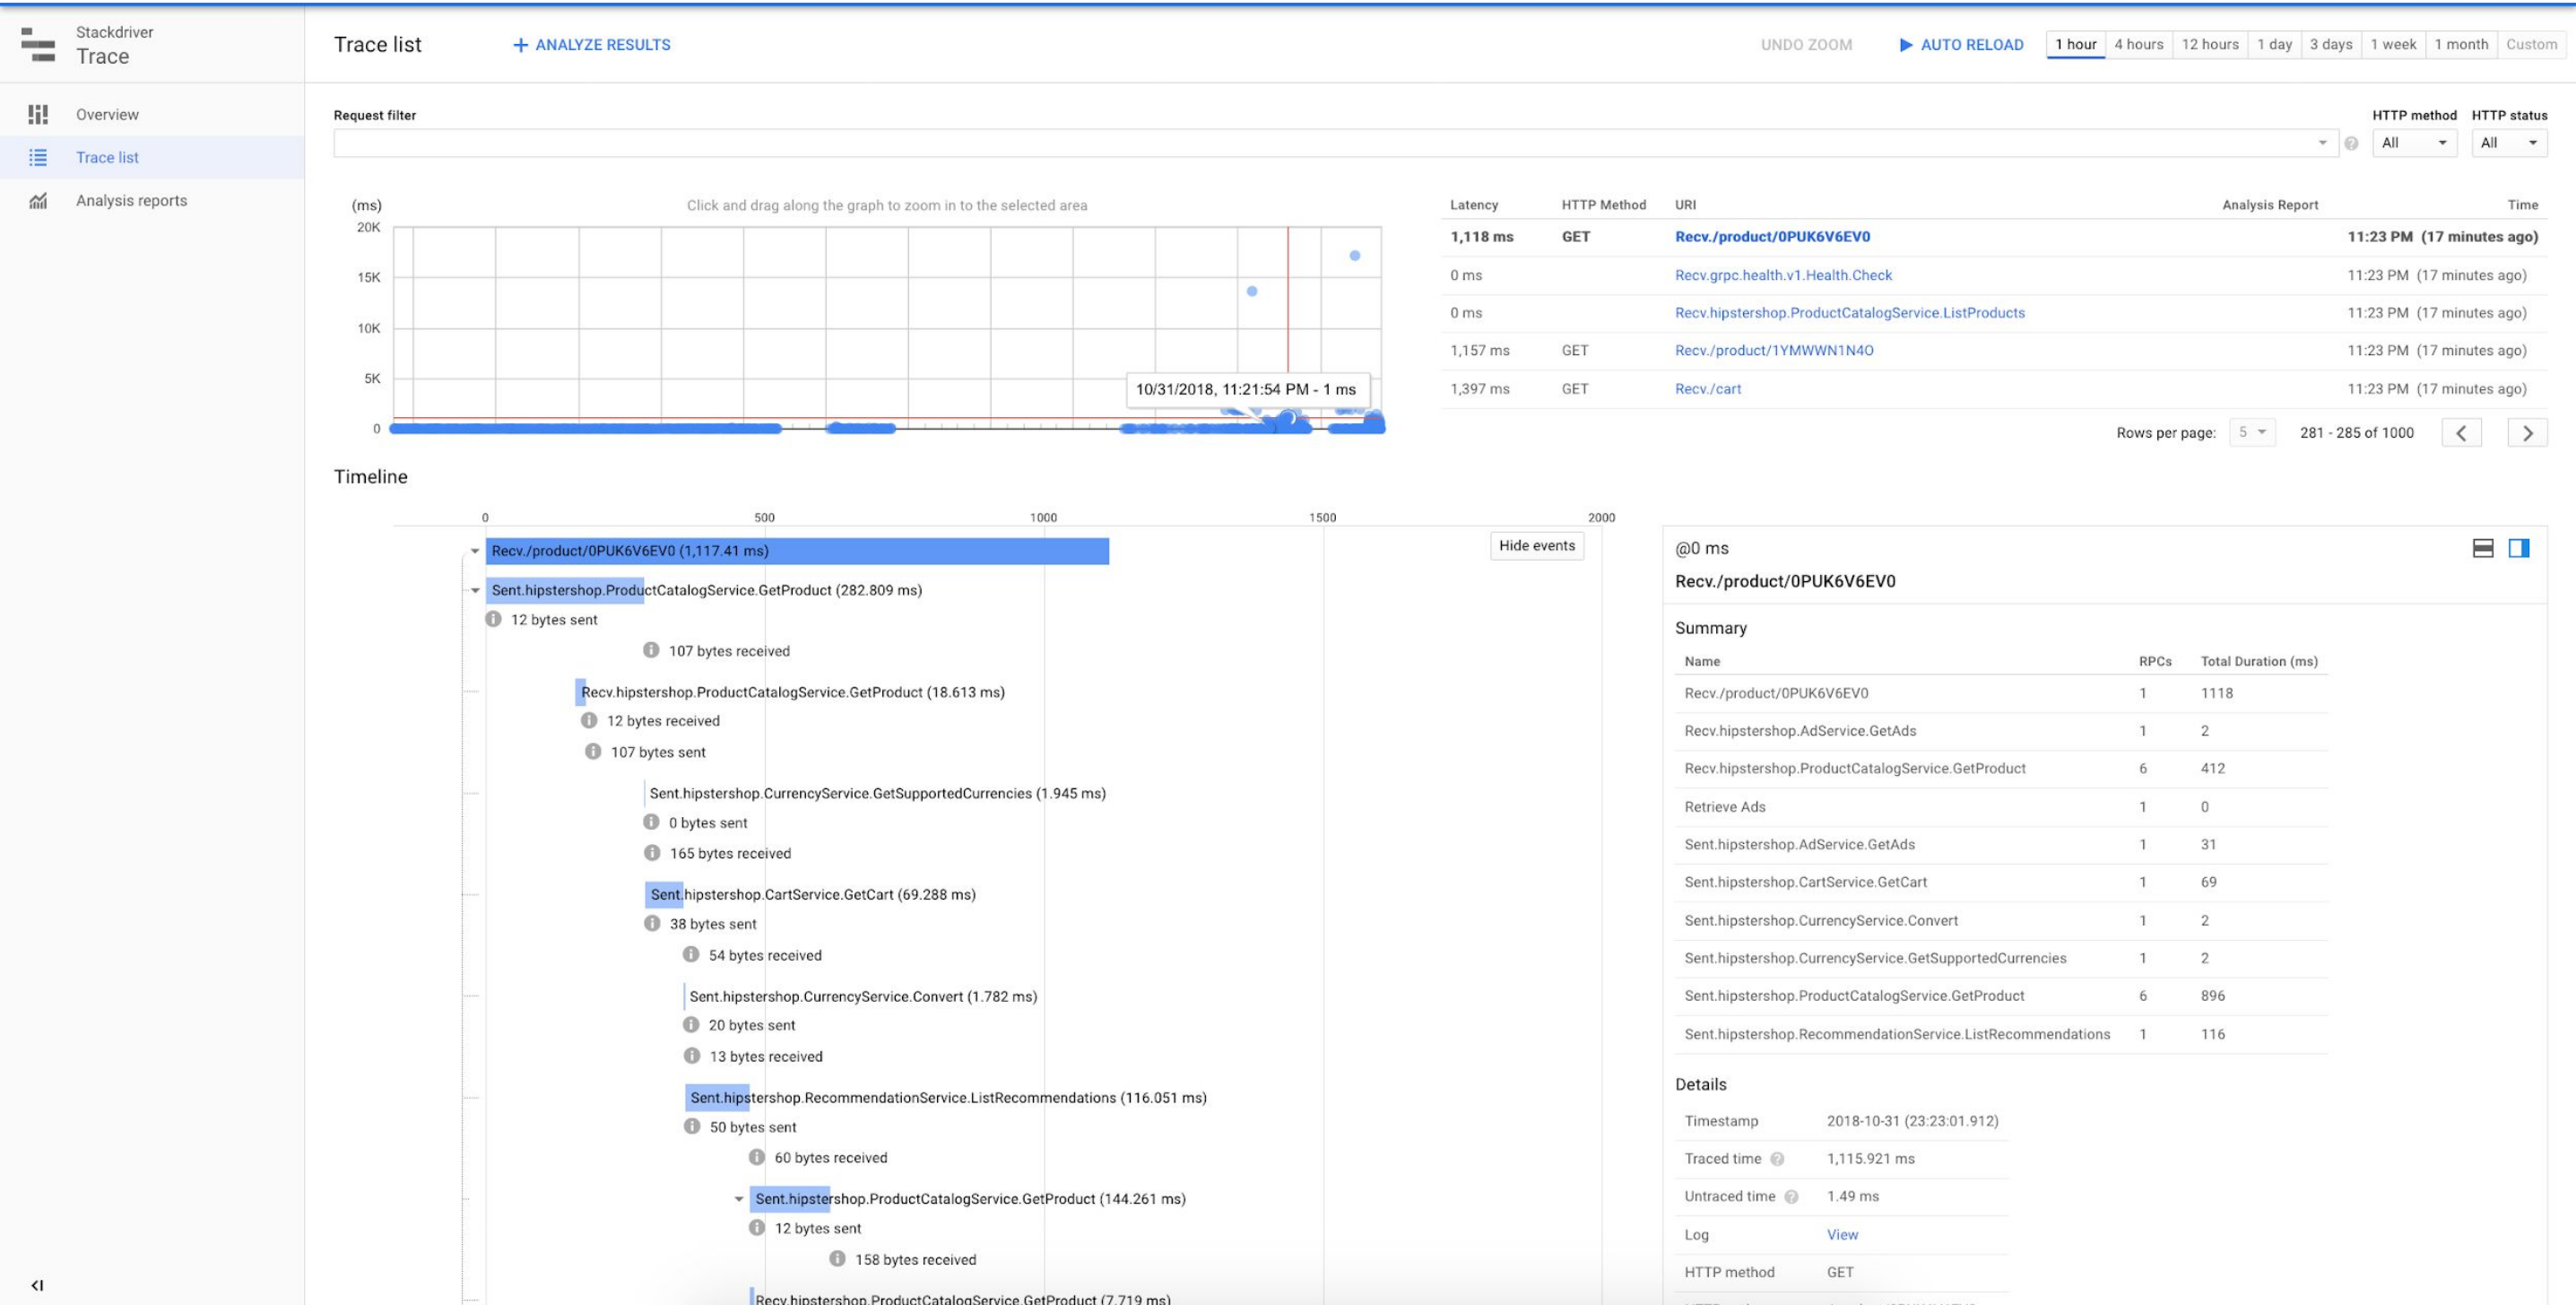
OpenCensus is an open source library for distributed tracing and monitoring. There is overlap between the OpenCensus APIs and the Stackdriver APIs, so that an alternative would be to develop the probe with Stackdriver APIs. However, there are several advantages of using OpenCensus instead:

1. You avoid an explicit dependency on a proprietary API. This will be important to people developing open source projects and companies wanting vendor independence. See the [GitHub repository](#) for the code.
2. OpenCensus has additional features for metric aggregation, which can intelligently help reduce the volume of monitoring data and controlling monitoring costs. See the [Aggregation](#) type for details on how to customize this.

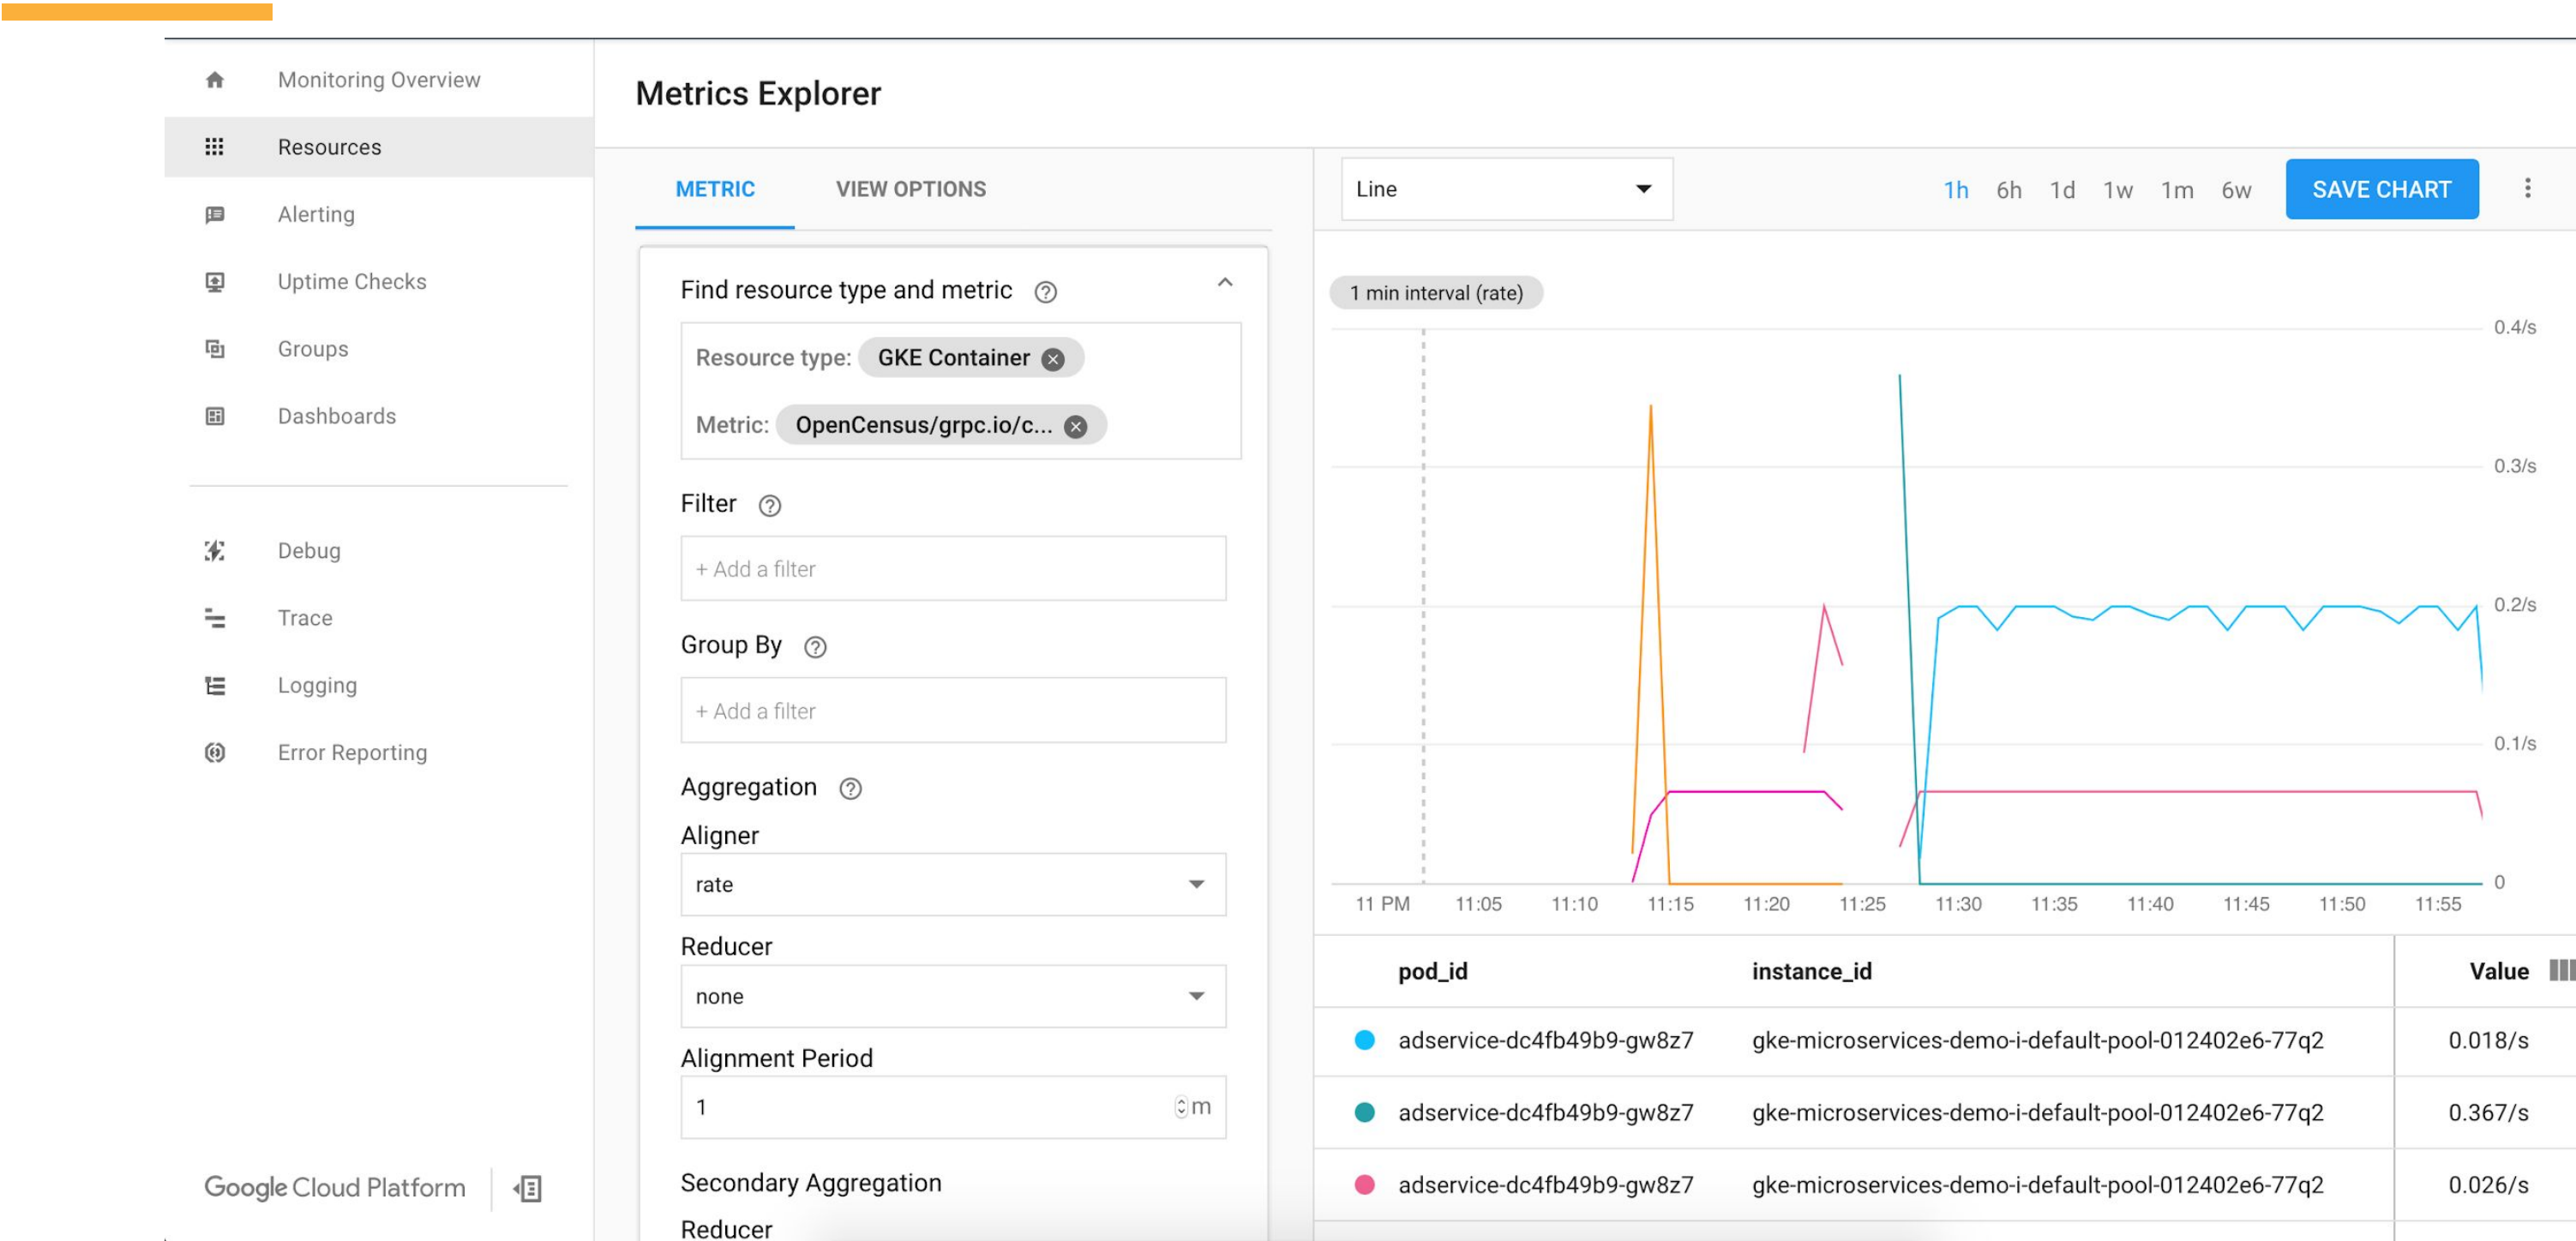
Distributed tracing systems: Opencensus

- 
3. OpenCensus supports exporters to multiple backends. That enables it development of code for collectors that can export monitoring stats to multiple clouds. This post uses the [Stackdriver exporter](#) to send the data to Google Cloud. See the [OpenCensus documentation](#) for other exporters.
 4. OpenCensus supports has prebuilt instrumentation for native library features, such as HTTP and gRPC. The code for this post uses the OpenCensus [HTTP plugin](#), which reduces the amount of code that needs to be written.

Distributed tracing systems: Opencensus



Distributed tracing systems: Opencensus



Distributed tracing systems: Conclusion

In the end, how important each piece is depends heavily on who you are and what kind of system you are building. For example, open source library authors are very interested in the OpenTracing API, while service developers tend to be more interested in the Trace-Context specification. When someone says one piece is more important than the other, they usually mean “one piece is more important to *me* than the other.”

However, the reality is this: Distributed Tracing has become a necessity for monitoring modern systems. In designing the building blocks for these systems, the age-old approach—“decouple where you can”—still holds true. Cleanly decoupled components are the best way to maintain flexibility and forwards-compatibility when building a system as cross-cutting as a distributed monitoring system.

Error tracking tools

We all want our application to run as smoothly as possible, but that's not always the case. Once the application is up and running, we need to know if and when errors or exceptions are thrown. That's why there are numerous error tracking tools in the market.

What is an error tracking tool? What can you actually do with it? What are the differences between the tools? Answers to these questions, along with some other relevant questions are waiting for you in the following post. Get ready to squash some bugs.

Error tracking tools

Deploying fast and frequently might introduce a errors, exceptions and bugs into our application, and it's best to identify them before your users do.

Current error tracking tools:

1. Sentry
2. Airbrake
3. BugSnag
4. OverOps
5. StackHunter
6. Rollbar
7. Raygun

Error tracking tools

Sentry a tiny bit of open source code grew to become a full-blown error monitoring tool, that identifies and debugs errors in production. You can use the insights to know whether you should spend time fixing a certain bug, or if you should roll back to a previous version.

The dashboard lets you see stack traces, with support for source maps, along with detecting each error's URL, parameters and session information. Each trace can be filtered with app, framework or raw error views.

Supported languages:


JavaScript, Node.js, Python, Go, PHP, Ruby, C#, Objective-C and even Swift

Error tracking tools

Features

You can catch errors in real-time as you deploy, and every error includes information about software, environment and the users themselves. If you believe your users are a key point at understanding errors, Sentry lets you prompt users for feedback whenever they encounter errors. That allows comparing their individual experience to the data you already have.

Error tracking tools



Captain Planet / Heart

ISSUES
EVENTS
OVERVIEW
USER FEEDBACK
RELEASES

Star

Subscribe

Settings

Unresolved Issues

Sort by: Last Seen

is:unresolved

		GRAPH:	24H	14D	EVENTS	USERS
Error	<div> <div>TypeError</div> <div>poll(..../sentry/scripts/views.js)</div> </div> <div> <div>Object [object Object] has no method 'updateFrom'</div> <div>HEART-1D 6 days ago — 4 months old 1</div> </div>				26	1
Error	<div> <div>★ javax.servlet.ServletException</div> <div>org.hsqldb.jdbc.Util in throwError</div> </div> <div> <div>Something bad happened</div> <div>HEART-1G 6 days ago — 4 months old</div> </div>				26	1
Error	<div> <div>script-src</div> <div>example.com</div> </div> <div> <div>Blocked 'script' from 'example.com'</div> <div>HEART-1K 6 days ago — 4 months old</div> </div>				26	1
Error	<div> <div>ZeroDivisionError</div> <div>bin/raven in <main></div> </div> <div> <div>divided by 0</div> <div>HEART-1H 6 days ago — 4 months old</div> </div>				26	1

GCP Stackdriver

Stackdriver Logging provides you with the ability to filter, search, and view logs from your cloud and open source application services. Allows you to define metrics based on log contents that are incorporated into dashboards and alerts.



Stackdriver

GCP Stackdriver: Features

1. Debugger
2. Alert
3. Error reporting
4. Tracing
5. Rapid discovery
6. Logging
7. Uptime monitoring
8. Dashboard
9. Integrations
10. Profiling
11. Smart defaults

Conclusion

Keeping in mind that you need to be proactive in your monitoring (i.e., prevent problems before they reach users), you need a solution that can test for performance issues in your pre-production environments and baseline the performance of your production environment. This means you'll want to opt for a synthetic monitoring solution over real user monitoring. Remember that with synthetic monitoring, you'll be developing behavioral scripts to simulate user actions and your solution will monitor them at specified intervals for performance, including availability, functionality, and response times.

Thank you for your attention!
Questions?

CONTACTS:

Email: stas.kolenkin@gmail.com

Skype: stas.kolenkin